

#### Announcements:

- Dan's office hours on Fri, Oct 18 cancelled.
- Don't discuss midterm until grades are released.
- No midterm grade ETA right now, but reminder that we have a clobber policy.

Lectures 21–22

# Pipelining

**CS 61C, Fall 2024 @ UC Berkeley**

Slides credit: Dan Garcia, Bora Nikolic, Peyrin Kao

Slides template credit: Josh Hug, Lisa Yan

# Measuring Performance

---

Lectures 21–22, CS 61C, Fall 2024

## Measuring Performance

### Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

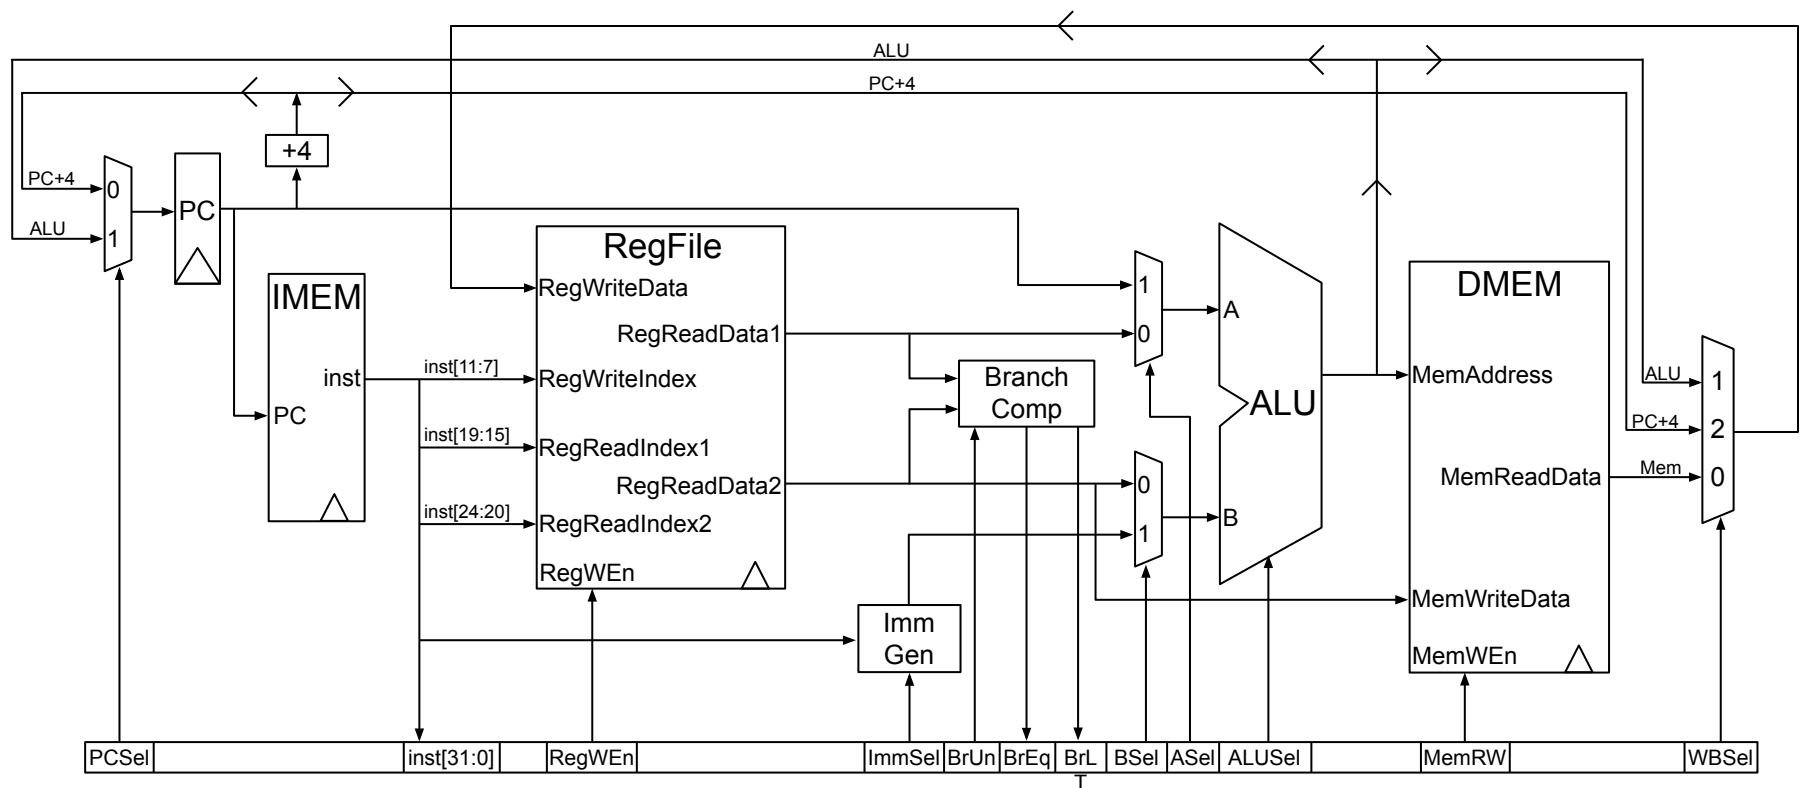
### Hazards

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

### Superscalar Processors

# The Single-Cycle RISC-V Datapath

Behold: Our circuit for running RISC-V instructions. We built this in the last 4 lectures.



# Instruction Timing

What is the clock period of our datapath circuit?

- Determined by longest path (lw).
- Clock period: 800ps.
- Frequency:  $1/800\text{ps} = 1.25 \text{ GHz}$ .
- 1.25 billion instructions per second.

Can we improve our datapath performance?

- What does it mean to improve performance?
- Quicker response time = one job finishes faster?
- More jobs per unit time?
- Longer battery life?

instr	Time taken per stage (in ps):					Total time taken:
	IF	ID	EX	MEM	WB	
add	200	100	200		100	600ps
beq	200	100	200			500ps
jal	200	100	200		100	600ps
<b>lw</b>	<b>200</b>	<b>100</b>	<b>200</b>	<b>200</b>	<b>100</b>	<b>800ps</b>
sw	200	100	200	200		700ps

This table shows some example timings.

Some cells are empty if the instruction doesn't use that stage.

# Analogy: Transportation

---

Consider two different vehicles:

	Sports Car	Bus
Passenger Capacity:	2	50
Travel Speed:	200 mph	50 mph
Gas Mileage:	5 mpg	2 mpg

For a 50 mile trip (assuming instant return trips)...

	Sports Car	Bus
Travel Time:	15 min	60 min
Time for 100 passengers:	750 min (50 trips)	120 min (2 trips)
Gallons per passenger:	5 gallons	0.5 gallons

# Measuring Performance in Computers

---

We'll look at 3 different measures of computer performance:

Transportation Analogy	In a Computer
Trip Time	<b>Program execution time</b> (Example: Time to update display)
Time for 100 passengers	<b>Throughput</b> (Example: Number of server requests handled per hour)
Gallons per passenger	<b>Energy per task</b> (Example: How many movies you can watch per battery charge)

## Optimizing Time Per Program

---

The "Iron Law" of processor performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \frac{\text{Cycles}}{\text{Instruction}} \frac{\text{Time}}{\text{Cycle}}$$

There are 3 components for optimizing the time it takes to execute a program.

The "Iron Law" of processor performance:

$$\frac{\text{Time}}{\text{Program}} = \boxed{\frac{\text{Instructions}}{\text{Program}}} \frac{\text{Cycles}}{\text{Instruction}} \frac{\text{Time}}{\text{Cycle}}$$

Instructions per program determined by:

- The task we're trying to do.
- The algorithm we implement, e.g.  $\theta(N^2)$  or  $\theta(N)$ .
- The programming language we use.
- The compiler we use.
- The instruction set architecture (ISA) we use, e.g. x86 or RISC-V.



The "Iron Law" of processor performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \boxed{\frac{\text{Cycles}}{\text{Instruction}}} \frac{\text{Time}}{\text{Cycle}}$$

Average clock cycles per instruction (CPI) determined by:

- The instruction set architecture (ISA).
- Processor implementation and microarchitecture, e.g. our datapath design.

Examples:

- For the datapath we made, CPI = 1.
- For complex instructions (e.g. strcpy), CPI > 1.
- We'll see CPI < 1 processors later.

The "Iron Law" of processor performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \frac{\text{Cycles}}{\text{Instruction}} \boxed{\frac{\text{Time}}{\text{Cycle}}}$$

Time per cycle (1/frequency) determined by:

- Processor microarchitecture: Critical path through logic gates.
- Technology: How many transistors fit in a given space.
- Power budget: Lower voltages reduce transistor speed.

# Optimizing Time Per Program

The "Iron Law" of processor performance:

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \frac{\text{Cycles}}{\text{Instruction}} \frac{\text{Time}}{\text{Cycle}}$$

	Processor A	Processor B
# Instructions	1 million	1.5 million
Average CPI	2.5	1
Clock rate <i>f</i>	2.5 GHz	2 GHz
Execution time	1 ms	0.75 ms

Processor B is faster for this task, despite executing more instructions and having a slower clock rate!

## Optimizing Energy Per Program (1/2)

---

There are 2 components for optimizing the energy it takes to execute a program.

$$\frac{\text{Energy}}{\text{Program}} = \boxed{\frac{\text{Instructions}}{\text{Program}}} \frac{\text{Energy}}{\text{Instruction}}$$

We already saw instructions per program earlier – it appears in this equation, too.

## Optimizing Energy Per Program (2/2)

---

There are 2 components for optimizing the energy it takes to execute a program.

$$\frac{\text{Energy}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \boxed{\frac{\text{Energy}}{\text{Instruction}}}$$

Energy per instruction =  $CV^2$

- $C$  = capacitance. Depends on hardware technology and processor features.
- $V$  = supply voltage.

To improve performance, we can reduce capacitance and voltage.

- Newer processors improve efficiency from Moore's law and lower supply voltage.
- Reducing supply voltage too far causes "leakage power" - transistor switches don't fully turn off.
- In recent years, it's gotten harder to reduce supply voltage.

The energy "Iron Law":

$$\frac{\text{Tasks}}{\text{Second}} = \frac{\text{Joules}}{\text{Second}} \frac{\text{Tasks}}{\text{Joule}}$$

Performance                      Power                      Energy Efficiency

Energy efficiency is a key metric in all computing devices.

- If system is constrained by power, we need better energy efficiency to get more performance at the same power.
- For energy-constrained systems, we need better energy efficiency to prolong battery life.

# Pipelining

## Analogy: Laundry

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

### Pipelining

- **Analogy: Laundry**
- Pipelining RISC-V Instructions
- Pipelined Datapath

Hazards

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors

## Pipelining Analogy: Laundry

---

Andrew, Bora, Caroline, Dan are doing laundry.

Each person needs to Wash, Dry, Fold, and Stash.

- **Wash** using the washer takes 30 mins.
- **Dry** using the dryer takes 30 mins.
- **Fold** using the table takes 30 mins.
- **Stash** using the drawers takes 30 mins.



# Laundry Timing Diagram

Example: From 2:00 to 2:30, Andrew is folding.



# Sequential Laundry

---

Naive approach: Andrew does everything, then Bora does everything, and so on.

- How long for everybody to finish laundry? 8 hours.
- How long for one person to do their laundry? 2 hours.

	1:00	1:30	2:00	2:30	3:00	3:30	4:00	4:30	5:00	5:30	6:00	6:30	7:00	7:30	8:00	8:30
Andrew	Wash	Dry	Fold	Stash												
Bora					Wash	Dry	Fold	Stash								
Caroline									Wash	Dry	Fold	Stash				
Dan													Wash	Dry	Fold	Stash

# Sequential Laundry

Can we do better than this?

Idea: When Andrew's drying, nobody is using the washer.

	1:00	1:30	2:00	2:30	3:00	3:30	4:00	4:30	5:00	5:30	6:00	6:30	7:00	7:30	8:00	8:30
Andrew	Wash	Dry	Fold	Stash												
Bora					Wash	Dry	Fold	Stash								
Caroline									Wash	Dry	Fold	Stash				
Dan													Wash	Dry	Fold	Stash

# Pipelined Laundry

---

Pipelined approach: While Andrew uses the Dryer, Bora can use the Washer.

- How long for everybody to finish laundry? 3.5 hours.
- How long for one person to do their laundry? 2 hours. (Same as before!)

	1:00	1:30	2:00	2:30	3:00	3:30	4:00
Andrew	Wash	Dry	Fold	Stash			
Bora		Wash	Dry	Fold	Stash		
Caroline			Wash	Dry	Fold	Stash	
Dan				Wash	Dry	Fold	Stash

## Pipelined Laundry

---

At any given moment in time, multiple people are in the system.

- Different people are using different resources at the same time.
- Example: At 2:30, Dan is using the Washer, Caroline is using the Dryer, and so on.

	1:00	1:30	2:00	2:30	3:00	3:30	4:00
Andrew	Wash	Dry	Fold	Stash			
Bora		Wash	Dry	Fold	Stash		
Caroline			Wash	Dry	Fold	Stash	
Dan				Wash	Dry	Fold	Stash

# Pipelined Laundry

---

Over time, one person uses each of the resources, in order.

	1:00	1:30	2:00	2:30	3:00	3:30	4:00
Andrew	Wash	Dry	Fold	Stash			
Bora		Wash	Dry	Fold	Stash		
Caroline			Wash	Dry	Fold	Stash	
Dan				Wash	Dry	Fold	Stash

# Filling and Draining the Pipeline

The ideal speedup for everyone to finish is 4x, since 4 things are happening at once.

Limited by:

- Startup costs.  
"Filling" the pipeline.
- Shutdown costs.  
"Draining" the pipeline.
- At the start and end, some resources are unused.

	1:00	1:30	2:00	2:30	3:00	3:30	4:00	4:30	5:00	5:30	6:00	6:30
Andrew	W	D	F	S								
Bora		W	D	F	S							
Caroline			W	D	F	S						
Dan				W	D	F	S					
Eddy					W	D	F	S				
Jero						W	D	F	S			
Myrah							W	D	F	S		
Rosalie								W	D	F	S	
Peyrin									W	D	F	S

# Pipelines Are Limited By Slowest Stage

What if Dry only took 20 mins, instead of 30?

- Bora finishes Dry at 2:20.
- But Andrew finishes Fold at 2:30. Bora still has to wait.

Pipeline rate is limited by the slowest pipeline stage.

	1:00	1:30	2:00	2:30	3:00	3:30	4:00
Andrew	Wash	Dry	Fold	Stash			
Bora		Wash	Dry	Fold	Stash		
Caroline			Wash	Dry	Fold	Stash	
Dan				Wash	Dry	Fold	Stash



## Pipelining Laundry: Observations

---

**Latency** (time it takes to finish a single task) is unchanged.

- It still took 2 hours for each person to do laundry.

**Throughput** (number of jobs finished per hour) increased.

- Sequential: 4 people took 8 hours = 0.5 people per hour.
- Pipelined: 4 people took 3.5 hours = 1.14 people per hour.

Maximum throughput speedup = number of stages.

- If 4 people were doing laundry at any given time, we'd have 4x speedup.
- Limited by cost of filling and draining pipeline at the start and end.

Pipeline rate is limited by the slowest pipeline stage.

- Must wait for the slowest stage to finish before moving on to the next stage.
- Balancing the length of each pipeline stage is good for speedup.

# Pipelining RISC-V Instructions

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

## Pipelining

- Analogy: Laundry
- **Pipelining RISC-V Instructions**
- Pipelined Datapath

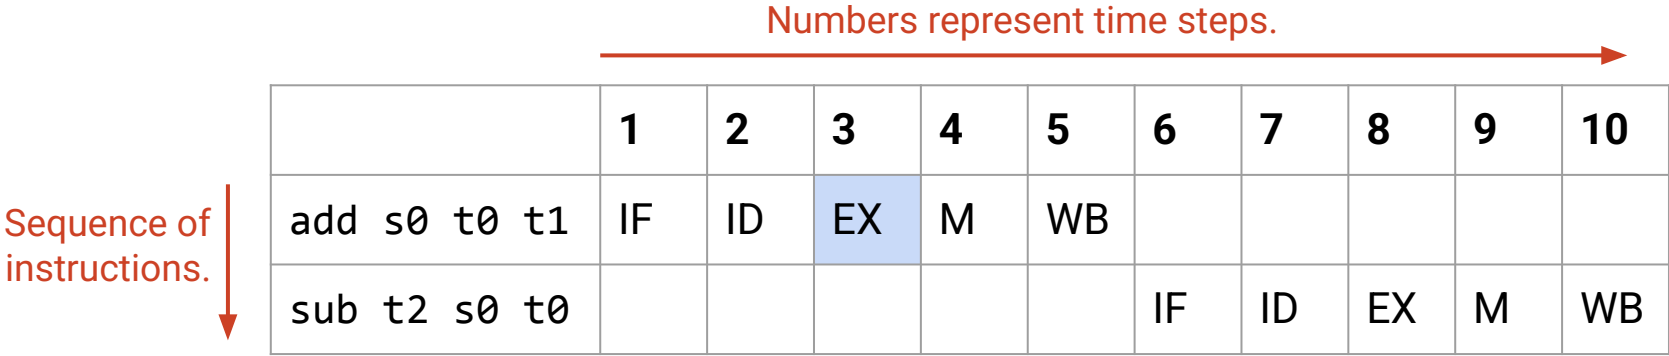
Hazards

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors

# Pipeline Timing Diagrams

Example: In the third time step, the add instruction is in the Execute (EX) stage of the pipeline.



# Sequential Instructions

---

Sequential approach: One instruction finishes, then the next one starts.

- How long to run 2 instructions? 10 time steps.
- How long to run each instruction? 5 time steps.

Our clock period would need to be 5 time steps.

	1	2	3	4	5	6	7	8	9	10
add s0 t0 t1	IF	ID	EX	M	WB					
sub t2 s0 t0						IF	ID	EX	M	WB

# Pipelined Instructions

---

Pipelined approach: Multiple instructions in the circuit at the same time.

- How long to run 2 instructions? 6 time steps.
- How long to run each instruction? 5 time steps. (Same as before!)

Our clock period is now 1 time step!

	1	2	3	4	5	6
add s0 t0 t1	IF	ID	EX	M	WB	
sub t2 s0 t0		IF	ID	EX	M	WB

# Pipelined Instructions

---

At any given moment in time, multiple instructions are in the circuit.

- Different instructions are using different resources at the same time.
- Example: At time step 3, the add instruction is using the Execute (EX) stage, while the sub instruction is using the Decode (ID) stage.

	1	2	3	4	5	6
add s0 t0 t1	IF	ID	EX	M	WB	
sub t2 s0 t0		IF	ID	EX	M	WB

# Pipelined Instructions

---

Over time, each instruction passes through all 5 pipeline stages.

- 1 stage per clock cycle.

Note: The pipelined CPU uses one clock for all stages.

- Clock cycle time is limited by the slowest stage (just like in the analogy).

	1	2	3	4	5	6
add s0 t0 t1	IF	ID	EX	M	WB	
sub t2 s0 t0		IF	ID	EX	M	WB

# Pipelined Datapath

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

## Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- **Pipelined Datapath**

Hazards

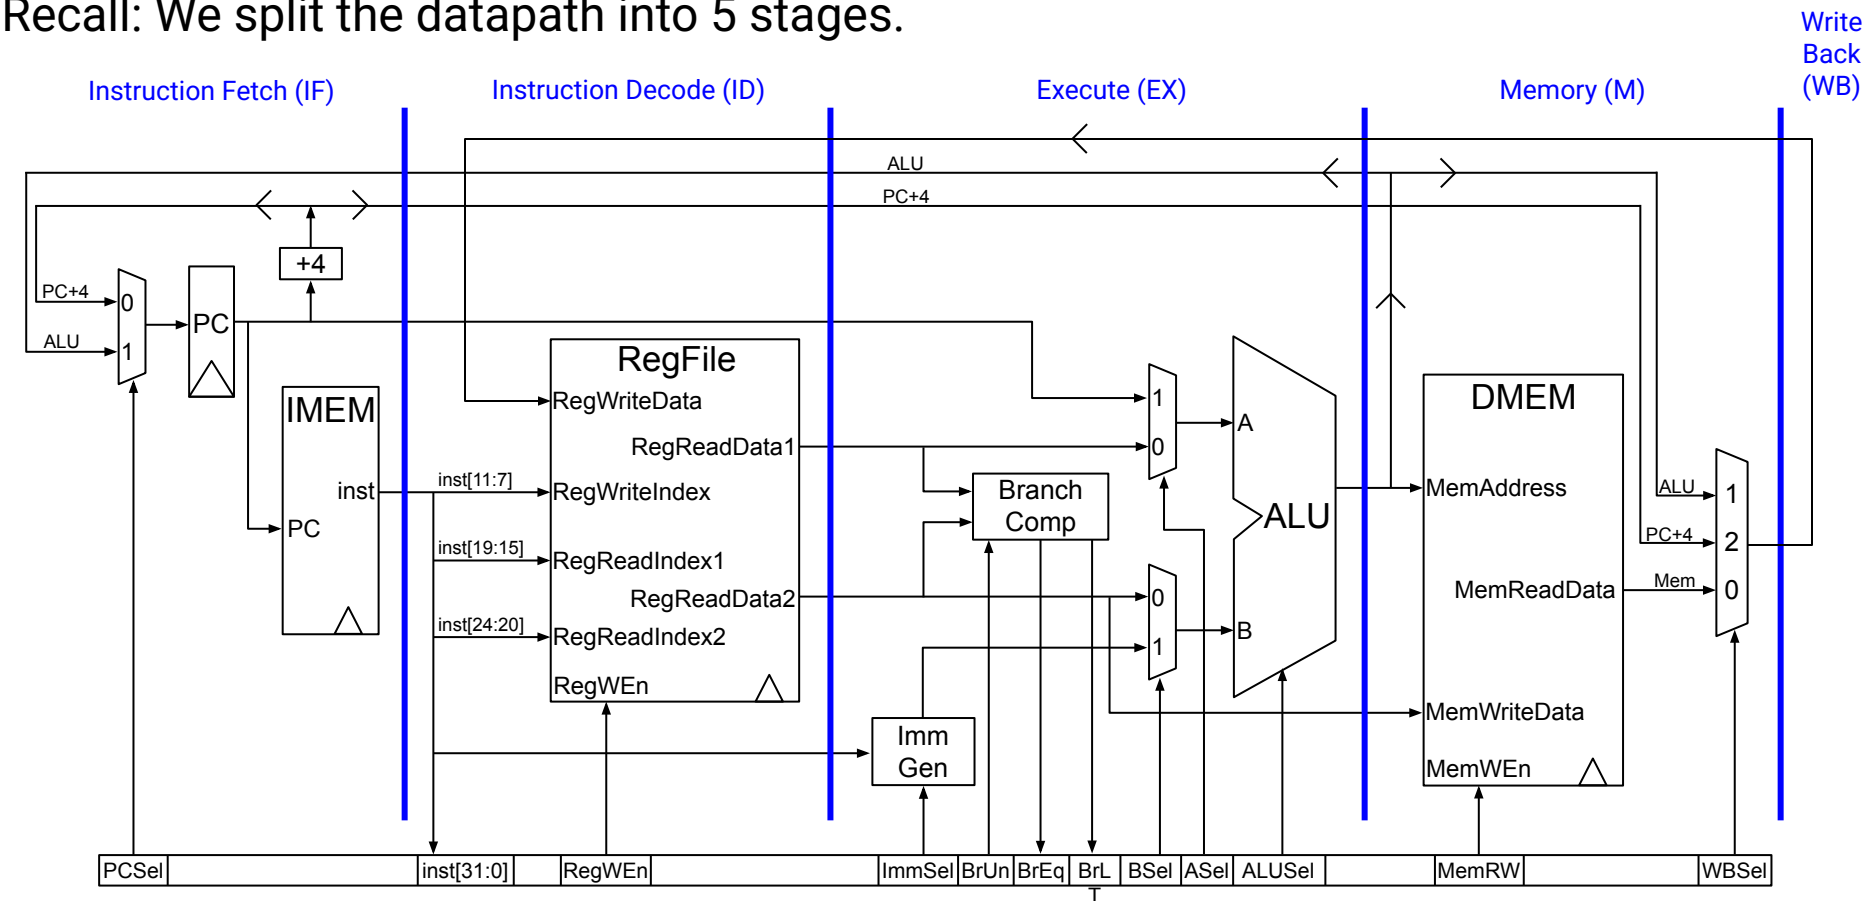
- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors



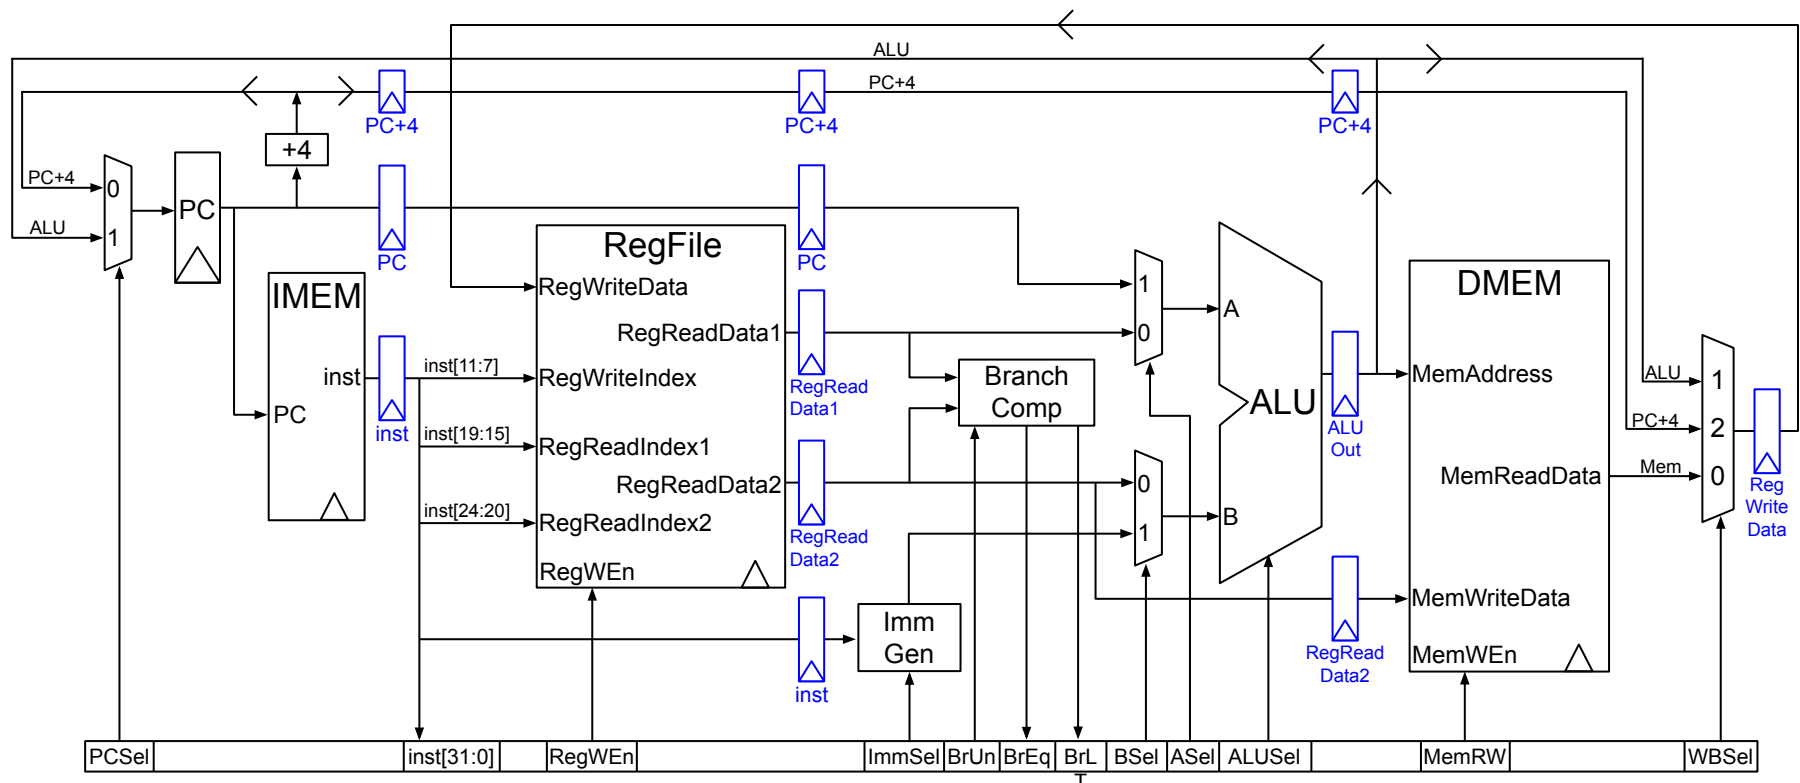
# Pipelined Datapath

Recall: We split the datapath into 5 stages.



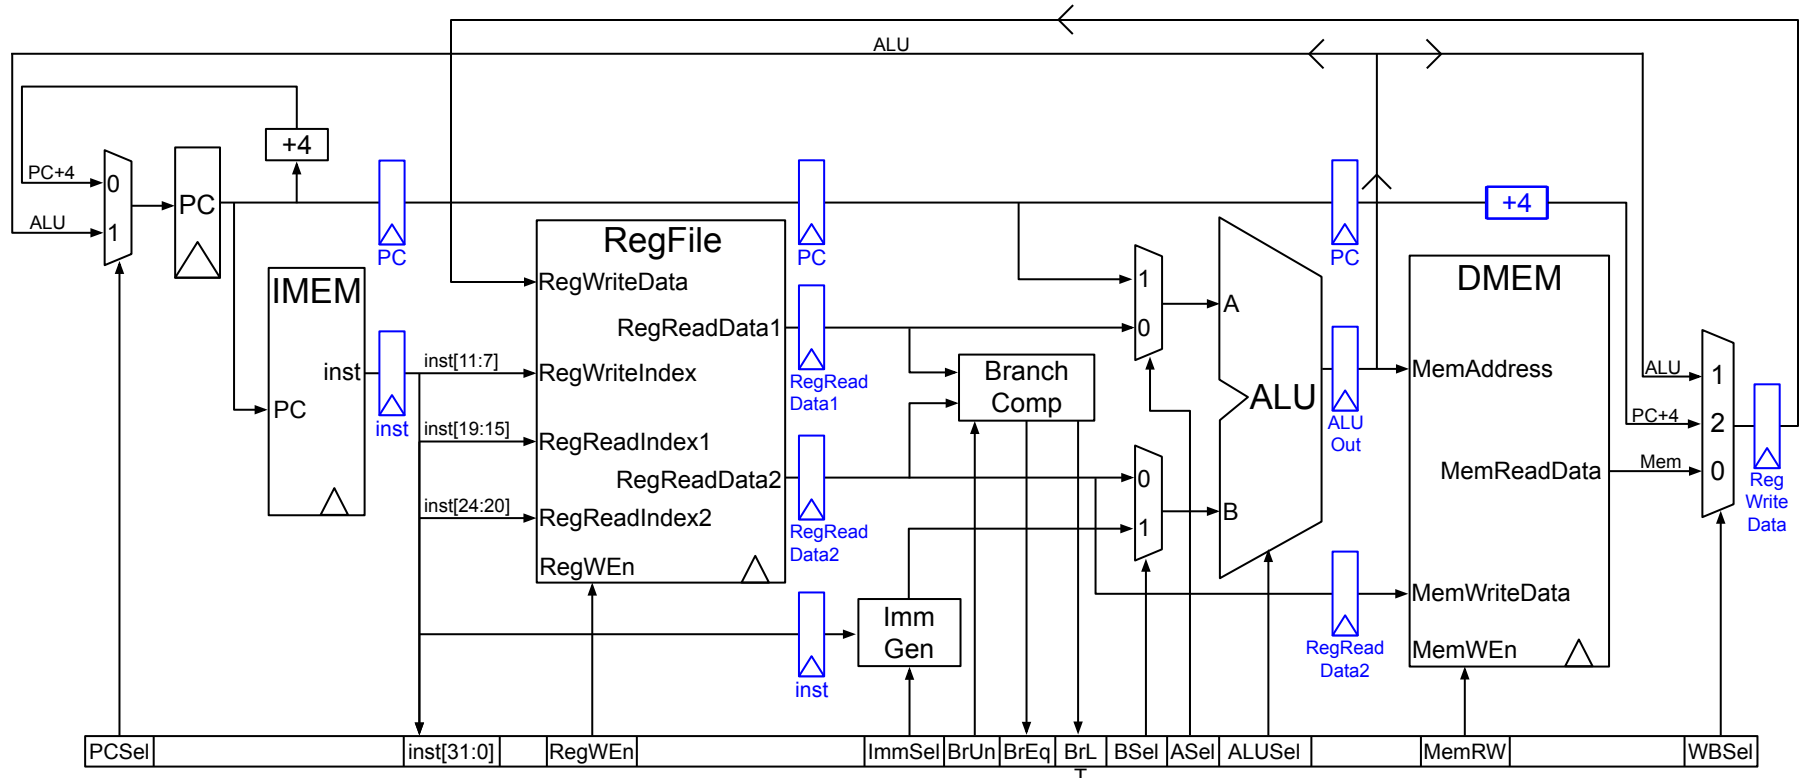
# Pipelined Datapath

Add registers between each stage to "hold" a signal until the next clock cycle.



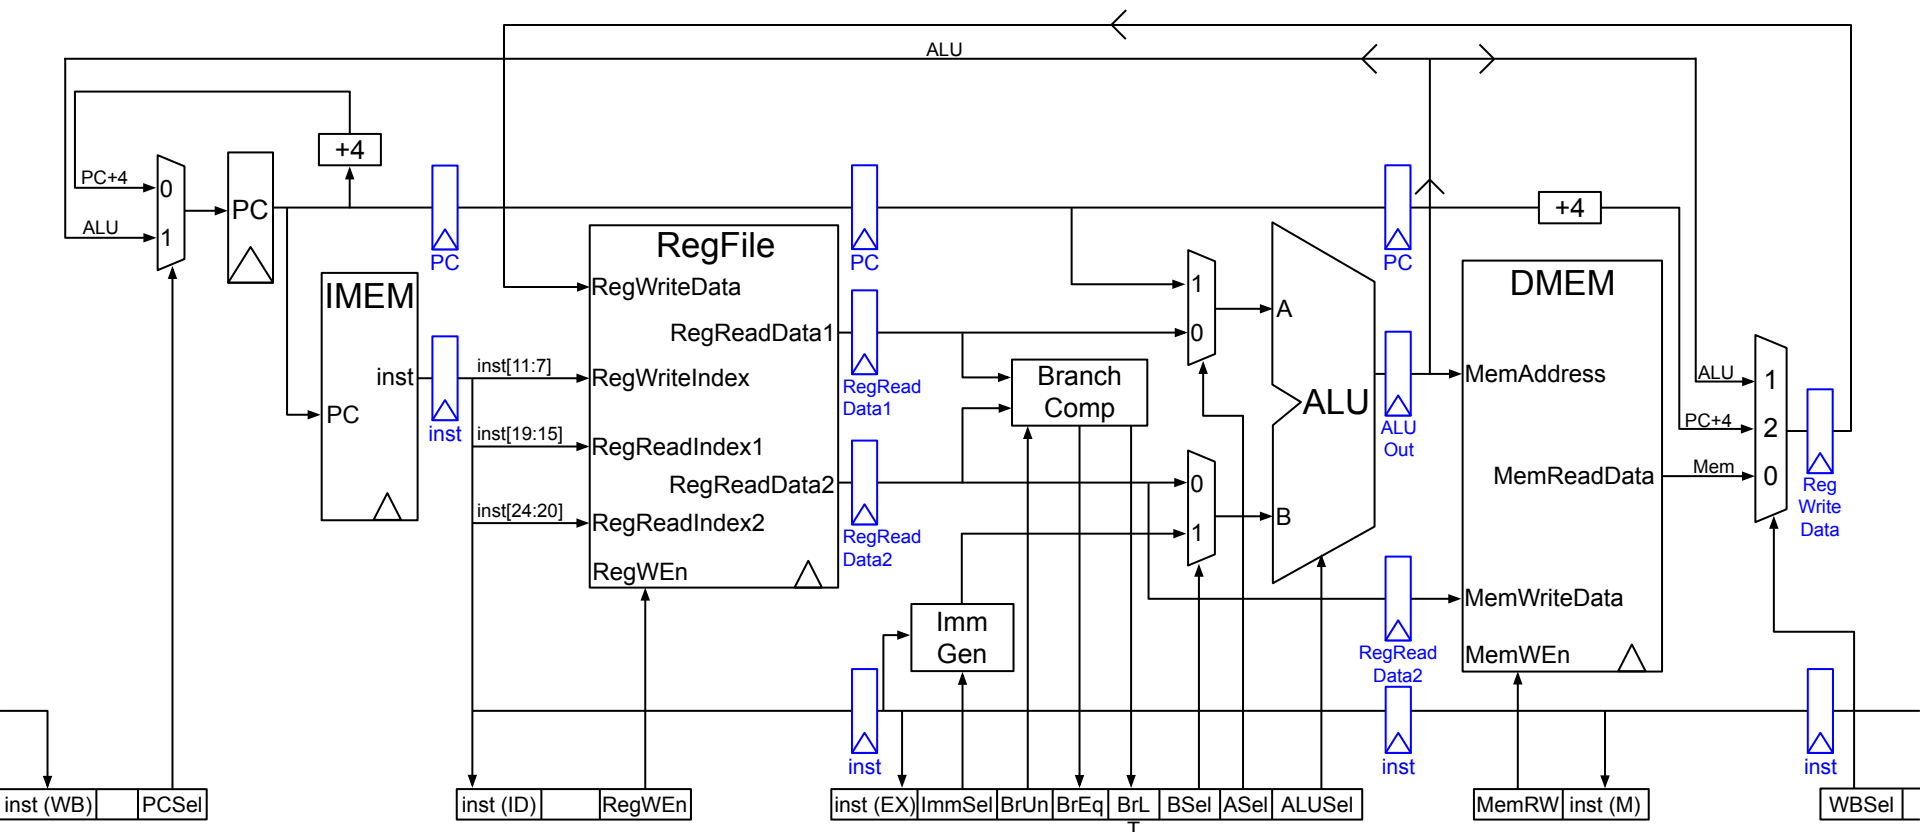
# Pipelined Datapath

Optimization: Recalculate PC+4 from PC in Memory stage to avoid storing/sending both PC and PC+4 down the pipeline.



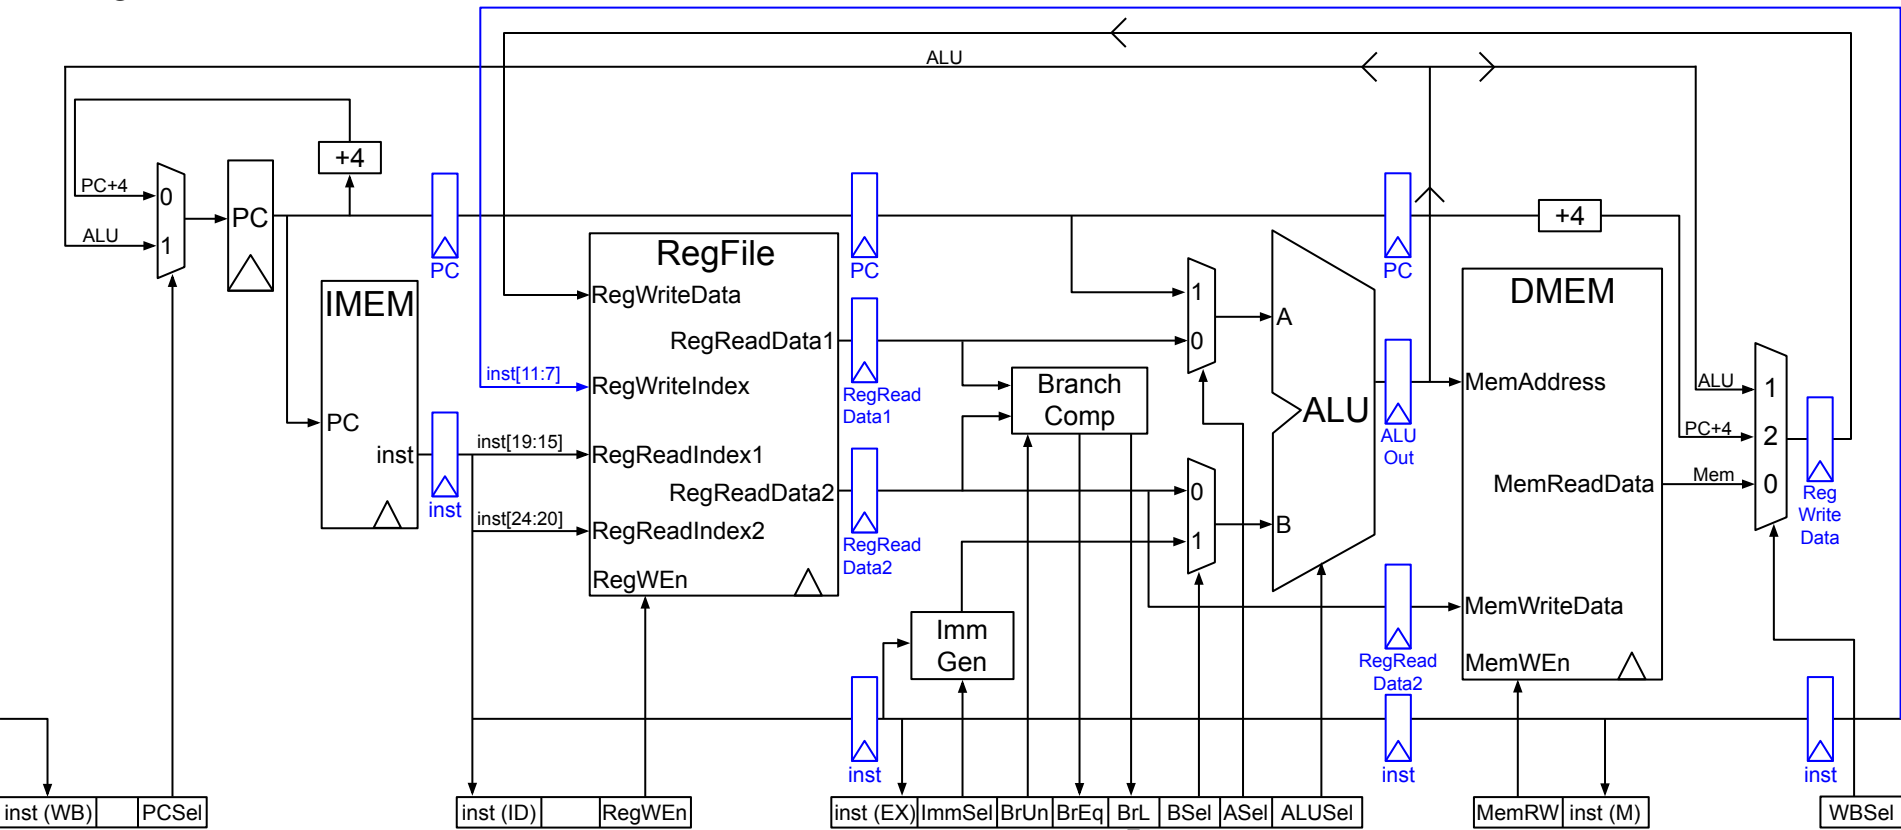
# Pipelined Datapath

Send instructions down the pipeline so control logic operates correctly in each stage.



# Pipelined Datapath

Change RegWriteIndex to use the rd field from the WB stage instruction, not the ID stage instruction.



Approach #1: One control logic subcircuit per stage.

- Store and send the instruction through the pipeline stages.
- Calculate relevant signals in each stage.
- Seen on the previous slides.

Approach #2: One control logic subcircuit.

- Calculate all the control logic signals in the decode (ID) stage.
- Send the control logic signals through the pipeline stages.

# Structural Hazards

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

**Hazards**

- **Structural Hazards**
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors

## Structural Hazards

---

Problem: Two or more instructions in the pipeline both need to access a single physical resource. In other words, a physical resource is needed in multiple stages.

- Analogy example: What if the "Fold" and "Stash" tasks both needed the table?

Solution 1: Instructions take turns to use the resource.

- Takes additional cycles per instruction.
- One instruction must *stall* (wait) while the other uses the resource.

Solution 2: Add more hardware to machine.

- Structural hazards can always be solved by adding more hardware.

Solution 3: Design the instruction set architecture (ISA) to avoid structural hazards.

- RISC-V is designed like this.
- Example: RegFile hardware can read two values per cycle. No RISC-V instruction needs to read 3 or more registers at once.



## Structural Hazard Example: Regfile

---

What are the physical limits of the Regfile?


- It can read from 2 registers in a single cycle.
- It can write to 1 register in a single cycle.

We avoid structural hazards by having separate "ports."

- 2 independent read ports, and 1 independent write port.
- Allows 3 simultaneous accesses per cycle.

Hypothetical structural hazards:


- What if the Regfile only had 1 read port?
- What if a RISC-V instruction needed to write to 2 registers?
- Good design helped us avoid structural hazards!



## Structural Hazard Example: Memory Access

The datapath has two memory blocks:

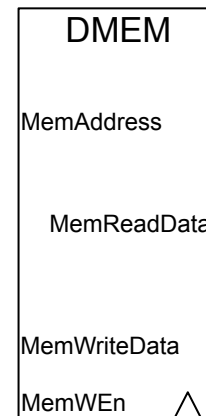
- IMEM: Instruction Memory.
- DMEM: Data Memory.
- But they're really the same memory.



In a pipelined datapath, IMEM and DMEM are used at the same time.

This is a structural hazard!

- Solution 1: IF stage *stalls* a cycle while MEM uses memory.
- Solution 2: Use two separate memories.
- Solution 3: Use caches to make IMEM and DMEM appear like separate memories.
  - This is what RISC-V does.
  - More on caches later!



# Data Hazards

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

**Hazards**

- Structural Hazards
- **Data Hazards**
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors

# Data Hazards

When does add **write** to register s0?      t = 5.

When does sub **read** the value in s0?      t = 3.

Is this a problem?

Reminder: Regfile lives in the ID stage, so that's why the read happens in the ID stage.

	1	2	3	4	5	6
add <b>s0</b> t0 t1	IF	ID	EX	M	<b>WB</b>	
sub t2 <b>s0</b> t0		IF	<b>ID</b>	EX	M	WB

# Data Hazards

Problem: An instruction depends on the result from the previous instruction.

- The previous instruction hasn't finished executing yet...
- ...so the current instruction doesn't know the result to use!

Example: When sub reads s0, add has not updated s0 yet.

	1	2	3	4	5	6
add <b>s0</b> t0 t1	IF	ID	EX	M	<b>WB</b>	
sub t2 <b>s0</b> t0		IF	<b>ID</b>	EX	M	WB

# Data Hazard Example (1/2)

---

Does this code have a data hazard?

	1	2	3	4	5	6	7
<code>add t0 t1 t2</code>	IF	ID	EX	M	WB		
<code>or t3 t4 t5</code>		IF	ID	EX	M	WB	
<code>slt t6 t0 t3</code>			IF	ID	EX	M	WB

## Data Hazard Example (1/2)

---

To identify data hazards: Check if we **write** to a register and then **read** from it later.

	1	2	3	4	5	6	7
<code>add t0 t1 t2</code>	IF	ID	EX	M	WB		
<code>or t3 t4 t5</code>		IF	ID	EX	M	WB	
<code>slt t6 t0 t3</code>			IF	ID	EX	M	WB

## Data Hazard Example (1/2)

---

To identify data hazards: Check if we write to a register and then read from it later.

- add **writes** to t0 at t = 5.
- slt **reads** from t0 at t = 4. This is too early to get the updated value!

	1	2	3	4	5	6	7
add t0 t1 t2	IF	ID	EX	M	<b>WB</b>		
or t3 t4 t5		IF	ID	EX	M	WB	
slt t6 t0 t3			IF	<b>ID</b>	EX	M	WB



## Data Hazard Example (2/2)

---

Does this code have a data hazard?

	1	2	3	4	5	6	7	8	9
<code>add t0 t1 t2</code>	IF	ID	EX	M	WB				
<code>or t3 t4 t5</code>		IF	ID	EX	M	WB			
<code>slt t6 t2 t3</code>			IF	ID	EX	M	WB		
<code>lw t3 4(t0)</code>				IF	ID	EX	M	WB	
<code>add t0 t1 t2</code>					IF	ID	EX	M	WB

## Data Hazard Example (2/2)

To identify data hazards: Check if we **write** to a register and then **read** from it later.

	1	2	3	4	5	6	7	8	9
<code>add t0 t1 t2</code>	IF	ID	EX	M	WB				
<code>or t3 t4 t5</code>		IF	ID	EX	M	WB			
<code>slt t6 t2 t3</code>			IF	ID	EX	M	WB		
<code>lw t3 4(t0)</code>				IF	ID	EX	M	WB	
<code>add t0 t1 t2</code>					IF	ID	EX	M	WB

## Data Hazard Example (2/2)

To identify data hazards: Check if we write to a register and then read from it later.

- add **writes** to t0 at t = 5.
- lw **reads** from t0 at t = 5. Is that too early?

	1	2	3	4	5	6	7	8	9
add <b>t0</b> t1 t2	IF	ID	EX	M	<b>WB</b>				
or t3 t4 t5		IF	ID	EX	M	WB			
slt t6 t2 t3			IF	ID	EX	M	WB		
lw t3 4( <b>t0</b> )				IF	<b>ID</b>	EX	M	WB	
add t0 t1 t2					IF	ID	EX	M	WB

## Data Hazard Example (2/2)

Some regfiles support writing a new value to a register, and then reading the new value, in the same cycle.

Not all regfiles support this. Check assumptions in any question.

	1	2	3	4	5	6	7	8	9
add <b>t0</b> t1 t2	IF	ID	EX	M	<b>WB</b>				
or t3 t4 t5		IF	ID	EX	M	WB			
slt t6 t2 t3			IF	ID	EX	M	WB		
lw t3 4( <b>t0</b> )				IF	<b>ID</b>	EX	M	WB	
add t0 t1 t2					IF	ID	EX	M	WB

# Fixing Data Hazards: Stalling

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

**Hazards**

- Structural Hazards
- Data Hazards
- **Fixing Data Hazards: Stalling**
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors

### Solution 1: **Stalling.**

- Wait for the first instruction to write its result before the second instruction reads the value.

### Solution 2: **Forwarding.**

- Add hardware to directly send the result back to the earlier stage, before the result is even written.

### Solution 3: **Code Scheduling.**

- Rearrange instructions to avoid data hazards.

Without stalling, we have a data hazard:

- add **writes** to t0 at t = 5.
- slt **reads** from t0 at t = 4 (too early to get the updated value).

We need slt to wait 2 time steps before reading t0.

	1	2	3	4	5	6	7
add t0 t1 t2	IF	ID	EX	M	<b>WB</b>		
or t3 t4 t5		IF	ID	EX	M	WB	
slt t6 t0 t3			IF	<b>ID</b>	EX	M	WB

# Stalling

Now, the read happens after the write. No more hazard!

- add **writes** to t0 at t = 5.
- slt **reads** from t0 at t = 6. The correct updated value is read!

	1	2	3	4	5	6	7	8	9
add t0 t1 t2	IF	ID	EX	M	<b>WB</b>				
nop		IF	ID	EX	M	WB			
nop			IF	ID	EX	M	WB		
or t3 t4 t5				IF	ID	EX	M	WB	
slt t6 t0 t3					IF	<b>ID</b>	EX	M	WB



To stall, we add no-ops into the code to delay an instruction.

- Recall: A no-op is an instruction that does nothing.
  - Example: `addi x0 x0 0`.
  - `nop` is a pseudo-instruction you can type in RISC-V.
- Compiler can insert no-ops to avoid data hazards.
  - Requires the compiler to know about the pipeline structure.

Stalls reduce performance, but they're required to get correct results!

# Fixing Data Hazards: Forwarding

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

**Hazards**

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- **Fixing Data Hazards: Forwarding**
- Fixing Data Hazards: Scheduling
- Control Hazards

Superscalar Processors

## Forwarding Example (1/2)

---

Problem: add writes t0 at t = 5, which is too late.

Idea: When do we first learn the result of the add computation?

	1	2	3	4	5	6
add t0 t1 t2	IF	ID	EX	M	WB	
slt t6 t0 t3		IF	ID	EX	M	WB

## Forwarding Example (1/2)

Problem: add writes t0 at t = 5, which is too late.

Idea: When do we first learn the result of the add computation?

- At t = 3, we already know the value that will be written to t0.

We can add hardware to *forward* this value so that slt can use it directly.

	1	2	3	4	5	6
add t0 t1 t2	IF	ID	EX	M	WB	
slt t6 t0 t3		IF	ID	EX	M	WB



Idea: Use the result as soon as it's computed.

- Read the value directly from a wire in the datapath.
- Don't wait for the result to be written to a register (and then read back).

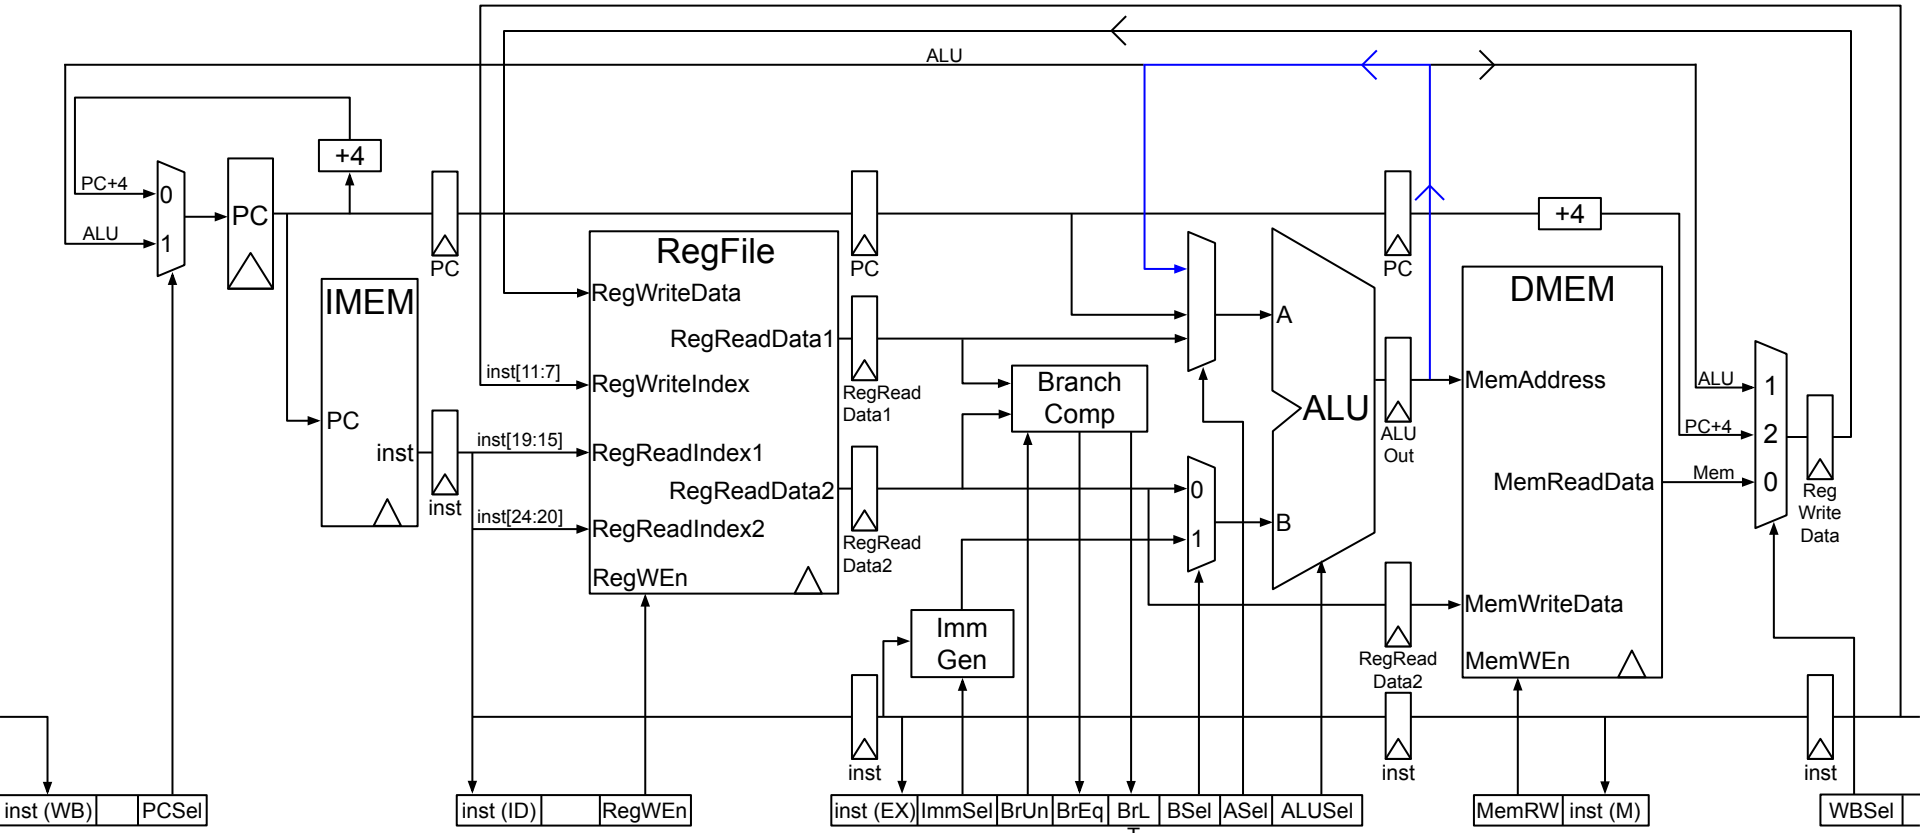
Forwarding requires extra connections in the datapath.

- We need extra wires to forward values to earlier stages.
- We need extra logic to check if forwarding is needed.
  - Where is the current instruction writing to? (i.e. what is rd?)
  - Where is the next instruction reading from? (i.e. what is rs1/rs2?)
  - If they're the same register\*, we need to forward.

\*Be careful to ignore writes to x0. If the current instruction is writing to x0, there's no data hazard, because writes to x0 do nothing.

# Adding Forwarding to Datapath (1/2)

This wire forwards the ALU result, so the next instruction can use result as ALU input.



## Forwarding Example (2/2)

---

Problem: lw writes to s2 at t = 5, which is too late.

Idea: When do we first learn the result of the load?

	1	2	3	4	5	6
lw <b>s2</b> 20(s1)	IF	ID	EX	M	<b>WB</b>	
add s4 <b>s2</b> s5		IF	<b>ID</b>	EX	M	WB

## Forwarding Example (2/2)

---

Problem: lw writes to s2 at t = 5, which is too late.

Idea: When do we first learn the result of the load?

- The value from memory shows up at t = 4.
- That's too late to send the value into the ALU!

We have to stall for one cycle.

	1	2	3	4	5	6
lw <b>s2</b> 20(s1)	IF	ID	EX	<b>M</b>	WB	
add s4 <b>s2</b> s5		IF	ID	<b>EX</b>	M	WB



## Forwarding Example (2/2)

Problem: lw writes to s2 at t = 5, which is too late.

Idea: When do we first learn the result of the load?

- The value from memory shows up at t = 4.
- That's too late to send the value into the ALU!

We have to stall for one cycle.

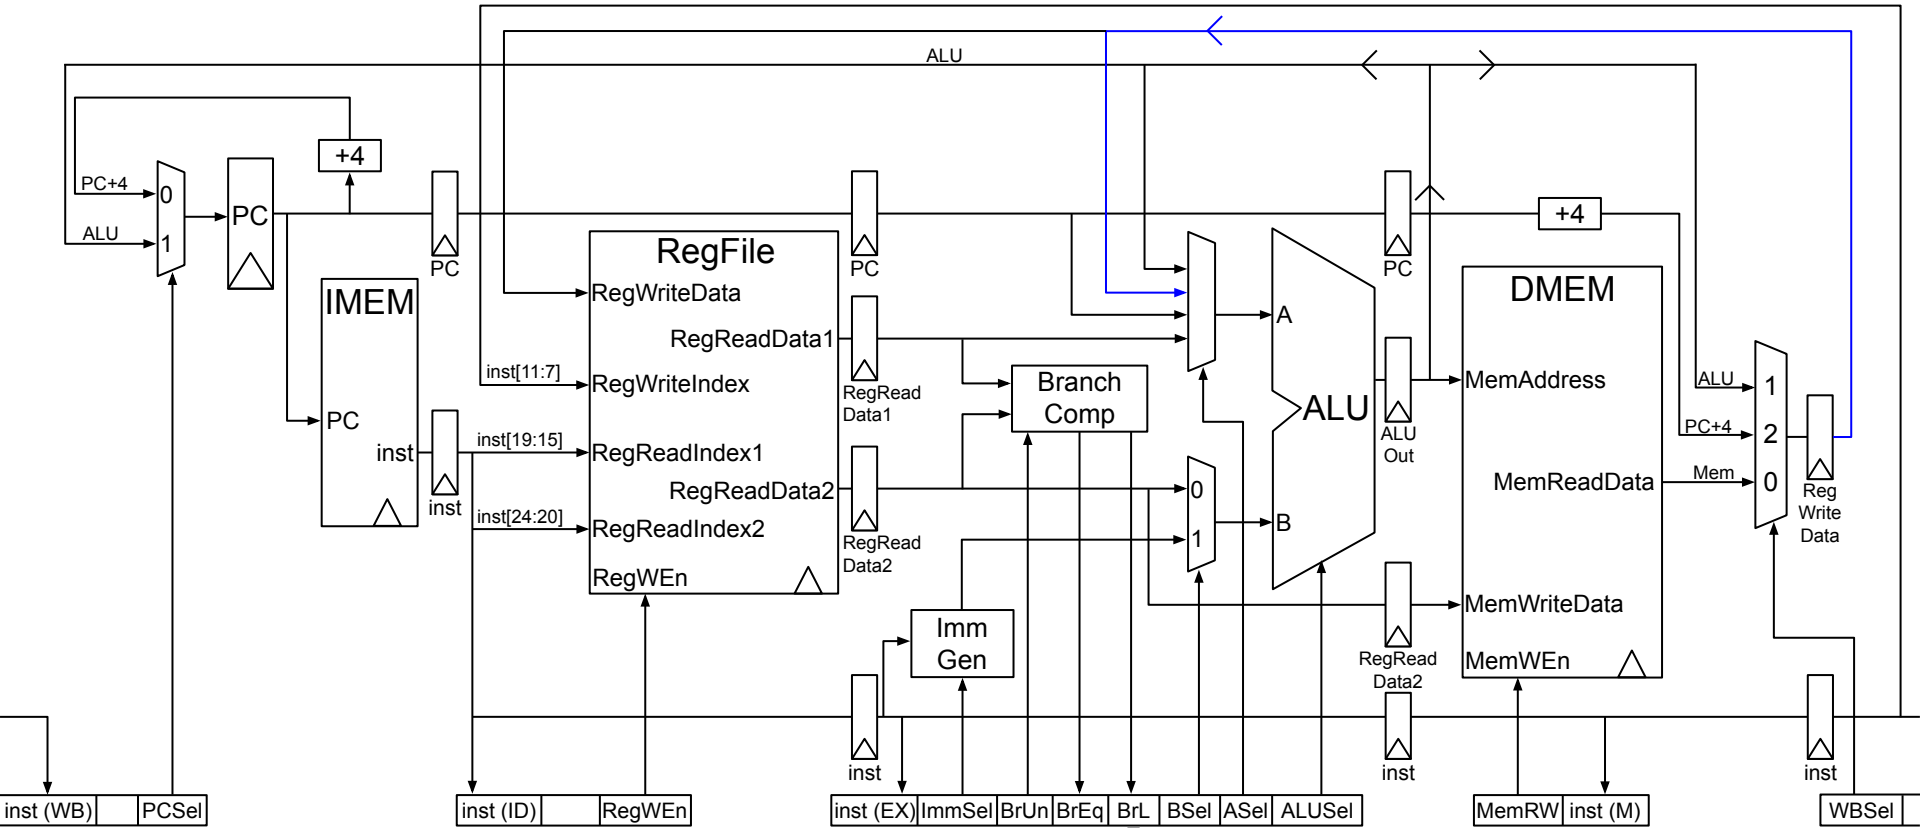
- Now, we can add hardware to forward the value so add can use it directly.

	1	2	3	4	5	6	7
lw s2 20(s1)	IF	ID	EX	M	WB		
nop		IF	ID	EX	M	WB	
add s4 s2 s5			IF	ID	EX	M	WB



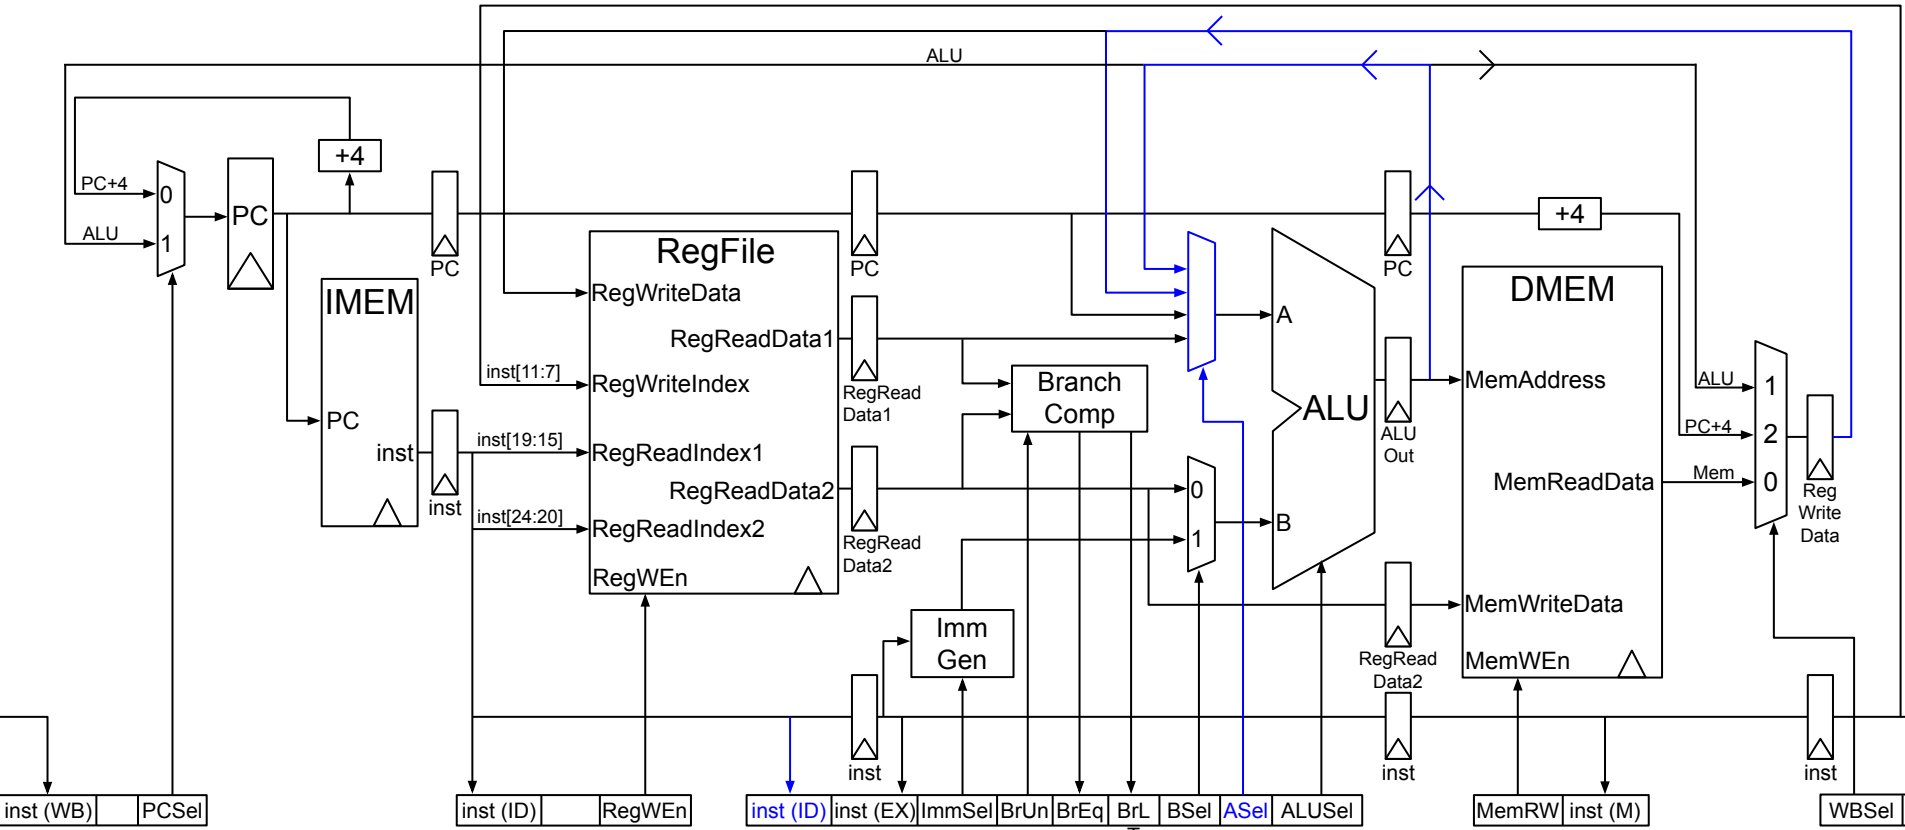
# Adding Forwarding to Datapath (2/2)

This wire forwards the value read from memory, so the next instruction can use the value as ALU input.



# Adding Forwarding Control Logic to Datapath

We need extra control logic to check whether we need to apply forwarding.



# Fixing Data Hazards: Code Scheduling

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

**Hazards**

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- **Fixing Data Hazards: Scheduling**
- Control Hazards

Superscalar Processors

Stalling inserts no-ops, which causes performance loss.

Idea: The compiler could use that no-op slot to execute an unrelated instruction.

- In other words, we *reorder* code to avoid data hazards.
- Requires the compiler to know about the pipeline structure.

# Code Scheduling Example

Consider this C code:     $A[3] = A[0] + A[1];$   
                               $A[4] = A[0] + A[2];$

Load A[0]	lw t0 0(a0)
Load A[1]	lw t1 4(a0)
A[0]+A[1]	add t3 t0 t1
Store A[3]	sw t3 12(a0)
Load A[2]	lw t2 8(a0)
A[0]+A[2]	add t4 t0 t2
Store A[4]	sw t4 16(a0)

Original instruction order: 2 data hazards.

Load A[0]	lw t0 0(a0)
Load A[1]	lw t1 4(a0)
Load A[2]	lw t2 8(a0)
A[0]+A[1]	add t3 t0 t1
Store A[3]	sw t3 12(a0)
A[0]+A[2]	add t4 t0 t2
Store A[4]	sw t4 16(a0)

After reordering the "Load A[2]" instruction:  
0 data hazards.

# Control Hazards

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

## Hazards

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- **Control Hazards**

Superscalar Processors

# Control Hazards

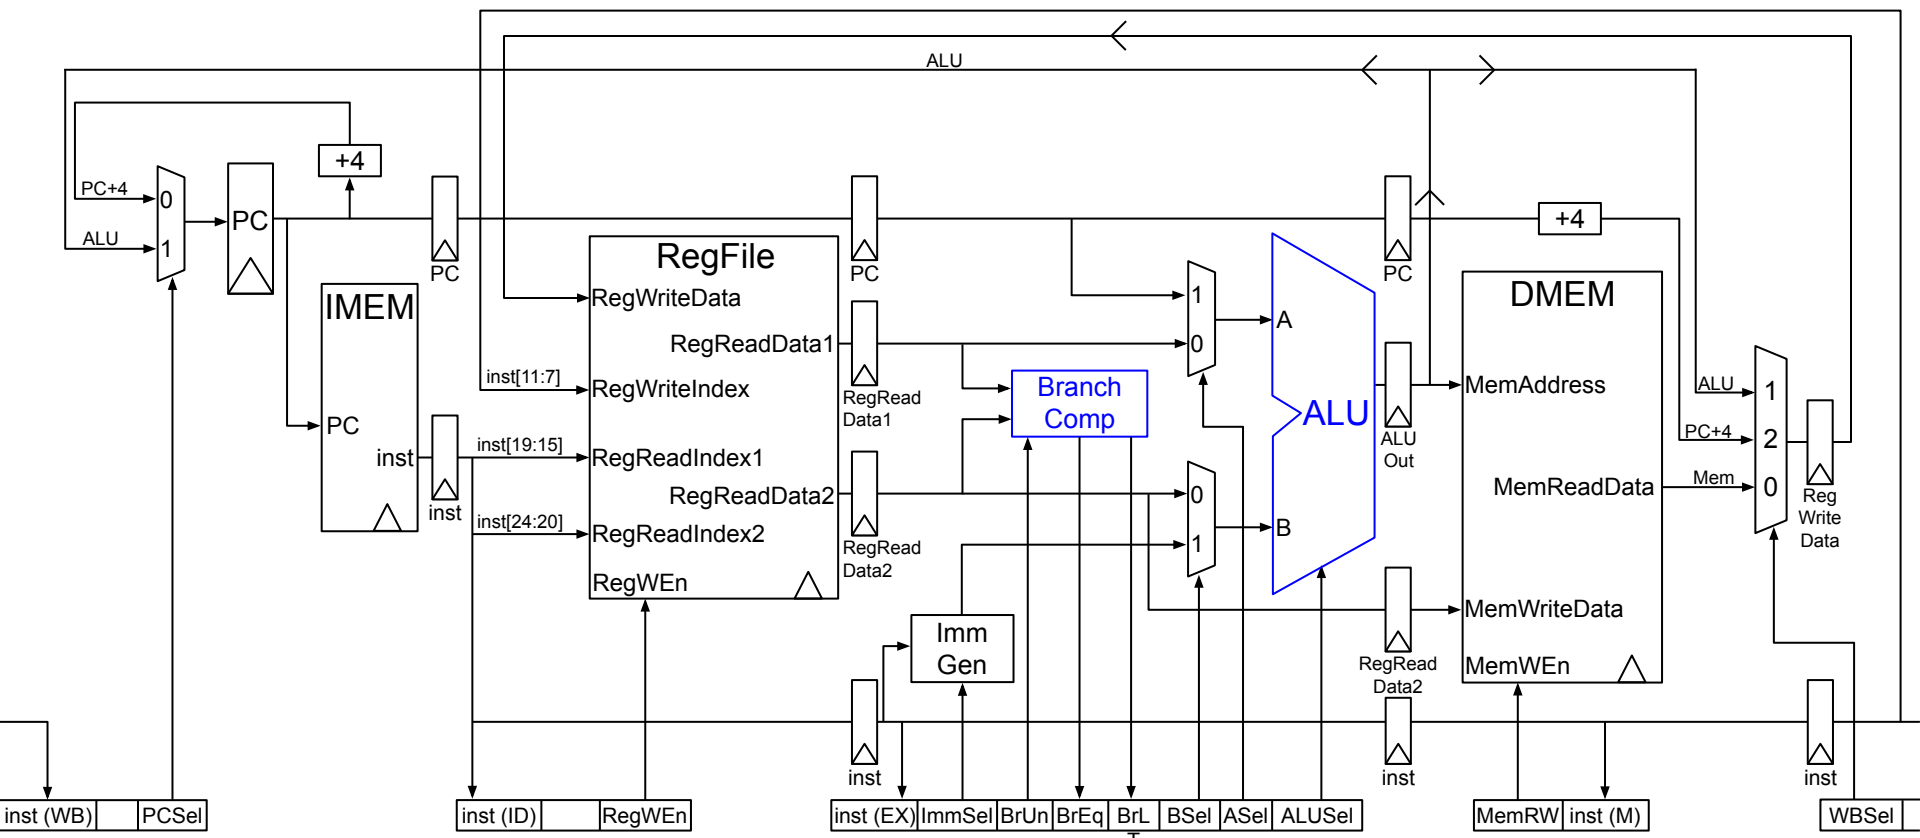
If the branch is taken ( $t0 == t1$ ), when do we figure out where to branch to?

	1	2	3	4	5	6	7	8	9
<code>beq t0 t1 Label</code>	IF	ID	EX	M	WB				
<code>sub t2 s0 t0</code>		IF	ID	EX	M	WB			
<code>add t6 s0 t3</code>			IF	ID	EX	M	WB		
<code>xor t5 t1 s0</code>				IF	ID	EX	M	WB	
<code>sw s0 8(t3)</code>					IF	ID	EX	M	WB



# Control Hazards

Recall: The EX stage is where we compute where to branch to.



## Control Hazards

If the branch is taken ( $t0 == t1$ ), when do we figure out where to branch to?

- At  $t = 3$ , the branch instruction runs the EX stage, which tells us where to go.
- But at  $t = 3$ , two more unwanted instructions have entered the pipeline!

	1	2	3	4	5	6	7	8	9
<code>beq t0 t1 Label</code>	IF	ID	EX	M	WB				
<code>sub t2 s0 t0</code>		IF	ID	EX	M	WB			
<code>add t6 s0 t3</code>			IF	ID	EX	M	WB		
<code>xor t5 t1 s0</code>				IF	ID	EX	M	WB	
<code>sw s0 8(t3)</code>					IF	ID	EX	M	WB

# Control Hazards

If the branch is taken, we're going somewhere else.

- We do not want to execute sub and add.
- We need to convert the next two instructions after the branch to no-ops.

	1	2	3	4	5	6	7	8	9
beq t0 t1 Label	IF	ID	EX	M	WB				
<del>sub t2 s0 t0</del>		<del>IF</del>	<del>ID</del>	<del>EX</del>	M	WB			
<del>add t6 s0 t3</del>			<del>IF</del>	<del>ID</del>	<del>EX</del>	M	WB		
xor t5 t1 s0				IF	ID	EX	M	WB	
sw s0 8(t3)					IF	ID	EX	M	WB

## Control Hazards

---

If a branch is not taken:

- The instructions fetched after the branch are correct.
- There's no control hazard.

If the branch is taken, or if there's a jump:

- The next two instructions in the pipeline are incorrect.
- There's a control hazard!
- To fix, we need to convert the incorrect instructions in the pipeline to no-ops.
- This fix is called *flushing* the pipeline.

## Branch Prediction

---

Problem: Every time we take a branch or jump, we lose 3 cycles to no-ops.

Solution: Branch prediction.

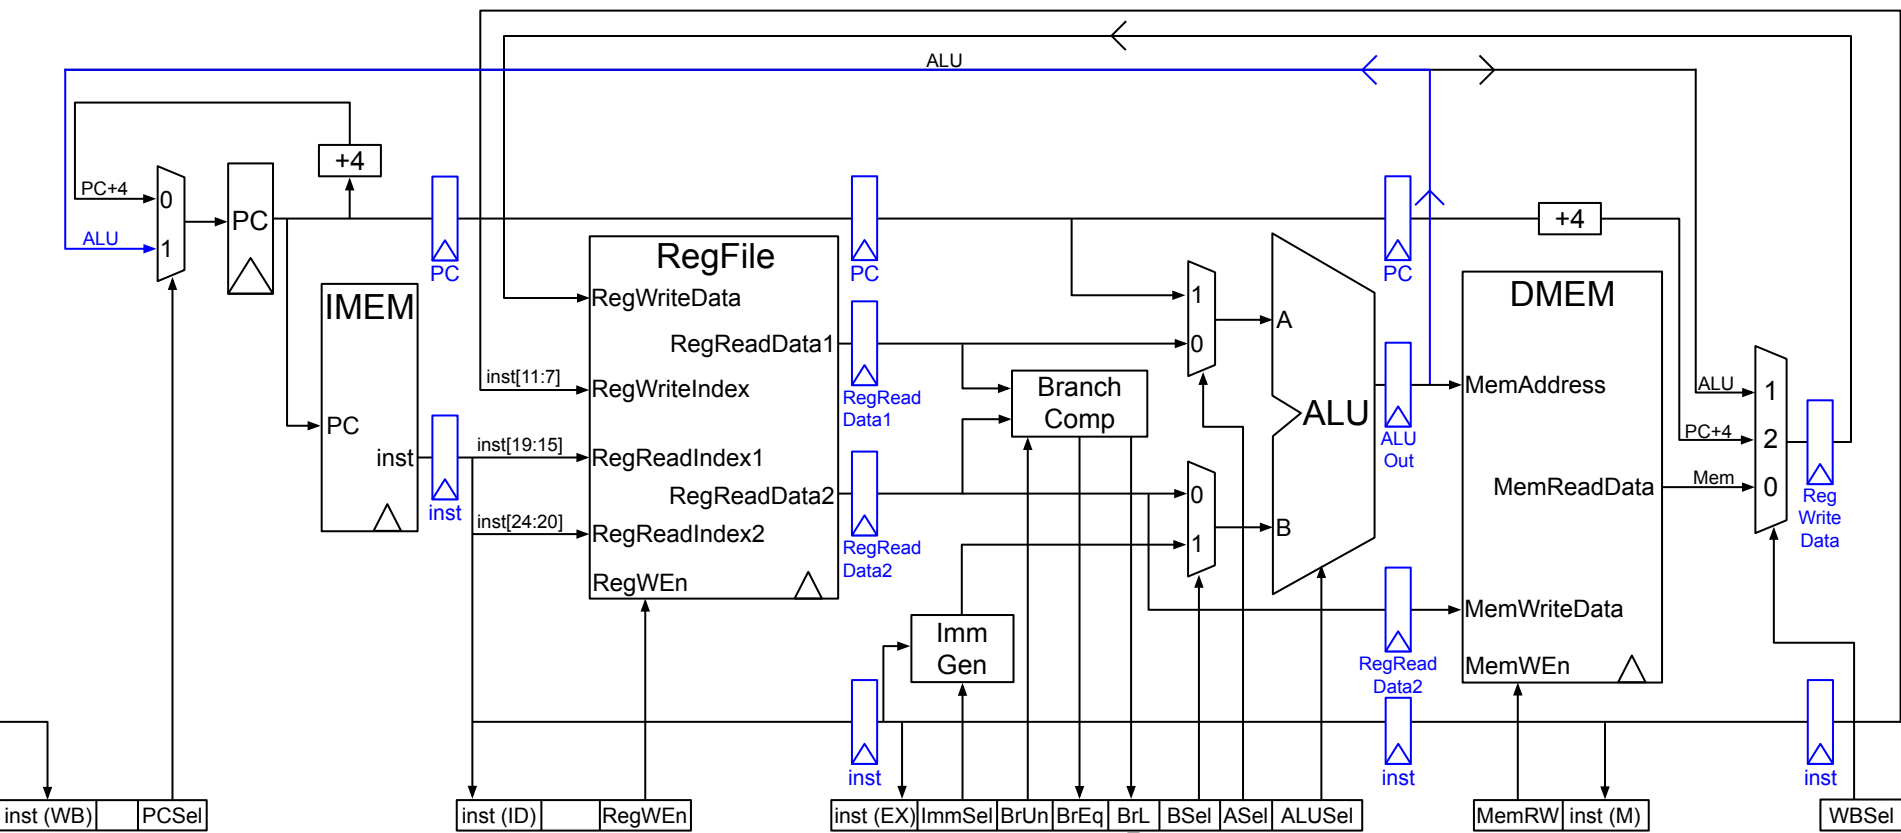
- Try and guess whether the branch will be taken.
- Only flush the pipeline if we guess wrong.
- Branch prediction improves performance *on average*, but there's still a penalty for wrong guesses.

What does our datapath do?

- Naive prediction: Our datapath always guesses "not taken" and fetches the next instructions.

# Pipelined Datapath

Note: We lose 3 cycles on a branch/jump, because updated PC is written back in MEM stage. While this happens, 3 instructions are in pipeline and need to be flushed.



# Superscalar Processors

---

Lectures 21–22, CS 61C, Fall 2024

Measuring Performance

Pipelining

- Analogy: Laundry
- Pipelining RISC-V Instructions
- Pipelined Datapath

Hazards

- Structural Hazards
- Data Hazards
- Fixing Data Hazards: Stalling
- Fixing Data Hazards: Forwarding
- Fixing Data Hazards: Scheduling
- Control Hazards

**Superscalar Processors**

The RISC-V ISA is designed for pipelining.

- All instructions are 32 bits wide.
  - Easy to fetch in one clock cycle, and decode in one clock cycle.
  - Contrast with CISC (R = Reduced, C = Complex) ISAs like x86:  
Instructions can be 1–15 bytes wide!
- There's a small set of standard instruction formats.
  - Can decode instruction and read from registers in one stage.
- Load/store addressing is conceptually simple.
  - In Stage 3 (EX), use the ALU to compute the address.
  - In Stage 4 (MEM), access memory.
- Memory operands are all aligned.
  - Memory access only takes one cycle.

The 5-stage pipeline is commonplace in many devices: Cars, appliances, etc.



## Further Increasing Processor Performance?

---

### 1. Increase clock rate.

- Limited by technology and power dissipation.

### 2. Increase pipeline depth.

- "Overlap" instruction execution through deeper pipeline, e.g. 10 or 15 stages.
- Less work per stage means shorter clock cycles and lower power.
- But more potential for all 3 types of hazards!
- More hazards means more stalling. Cycles per instruction  $> 1$ .

### 3. Design a *superscalar* processor.

- Desktops, laptops, cell phones, etc. often have a few of these, combined with simpler 5-stage pipeline processors.

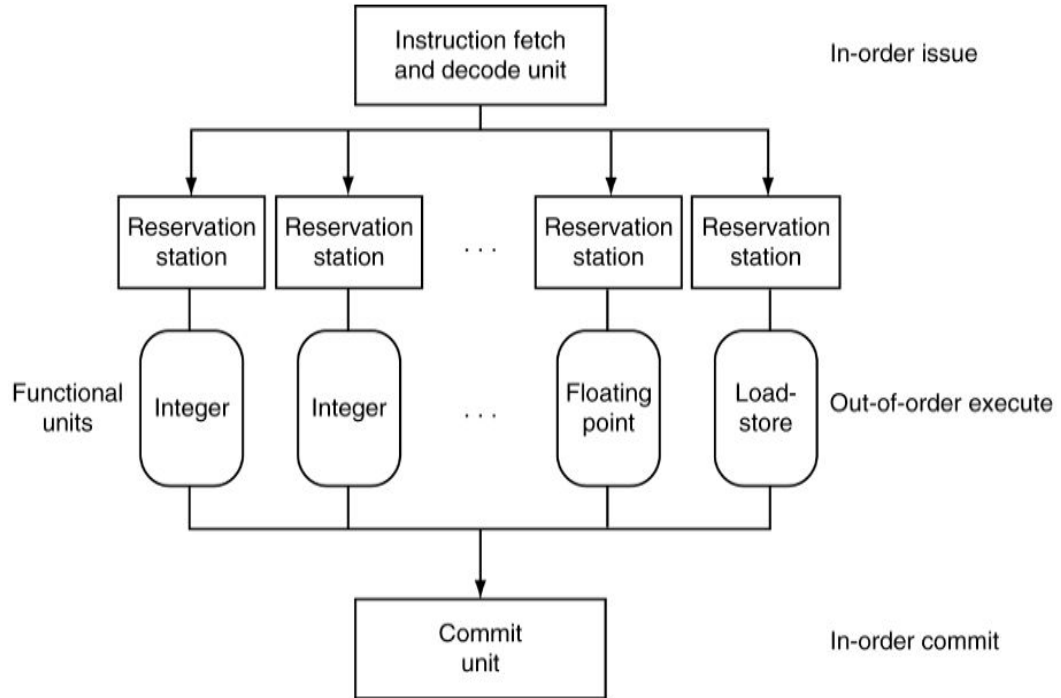
*Multiple-issue:* Start multiple instructions per clock cycle.

- Multiple execution units execute instructions in parallel.
- Each execution has its own pipeline.
- $CPI < 1$ : Multiple instructions completed per clock cycle.

*Dynamic "out-of-order" execution:*

- Reorder instructions dynamically in hardware to reduce impact of hazards.

Take CS 152 to learn more!



# Computing Cycles Per Instruction (CPI)

How do we measure CPI?

- In practice, we measure CPI on various benchmarks.
- There are benchmark programs from a variety of application domains:  
Data compression, code compilation, video decoding, network simulation, etc.

Recall the "Iron Law" of processor performance:

$$\underbrace{\frac{\text{Time}}{\text{Program}}}_{\text{Can measure.}} = \underbrace{\frac{\text{Instructions}}{\text{Program}}}_{\text{Can count.}} \boxed{\frac{\text{Cycles}}{\text{Instruction}}} \underbrace{\frac{\text{Time}}{\text{Cycle}}}_{1/\text{clock rate.}}$$

Rearrange terms to solve for CPI:

$$\boxed{\frac{\text{Cycles}}{\text{Instruction}}} = \frac{\text{Time}}{\text{Program}} \div \left( \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Time}}{\text{Cycle}} \right)$$

## ARM A53 Benchmark

ARM Cortex-A53 Core: 2 GHz,  
dual-issue processor:

- 4 BIPS (billion instructions per second).
- Peak CPI = 0.5.
- However, instruction/pipeline dependencies reduce performance.

If  $CPI < 1$ , we can also measure in terms of IPC (Instructions Per Cycle).

More about parallelism and performance later in CS 61C!

