



UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)

### Intro to RISC-V

# The RISC-V Instruction Set Architecture

- The RISC-V Instruction Set Architecture
- Elements of Architecture:  
Registers
- Add/Sub Instructions
- Immediates

# Great Idea #1: Abstraction (Levels of Representation/Interpretation)



**High Level Language  
Program (e.g., C)**

```
temp = v[k];
v[k] = v[k+1];
v[k+1] = temp;
```

*| Compiler*  
**Assembly Language  
Program (e.g., RISC-V)**

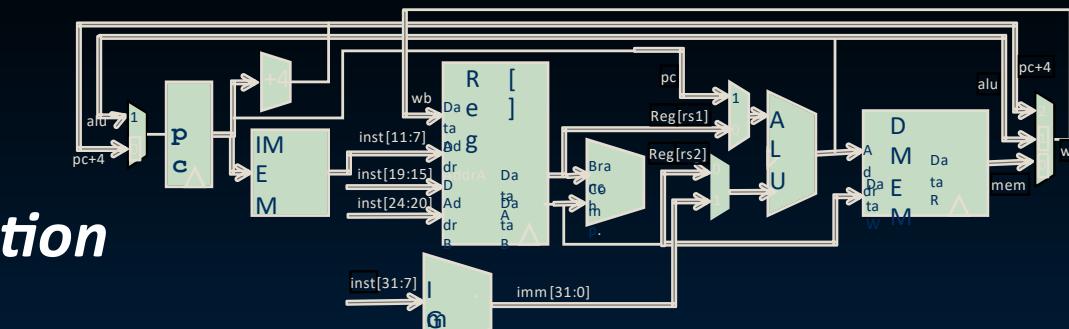
```
lw    x3, 0(x10)
lw    x4, 4(x10)
sw    x4, 0(x10)
sw    x3, 4(x10)
```

Anything can be represented  
as a number,  
i.e., data or instructions

*| Assembler*  
**Machine Language  
Program (RISC-V)**

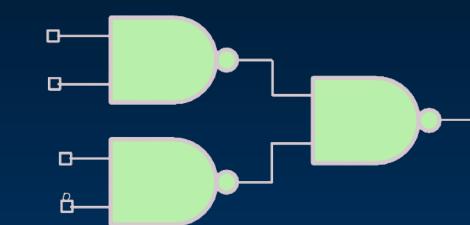
1000 1101 1110 0010 0000 0000 0000 0000 0000
1000 1110 0001 0000 0000 0000 0000 0000 0100
1010 1110 0001 0010 0000 0000 0000 0000 0000
1010 1101 1110 0010 0000 0000 0000 0000 0100

**Hardware Architecture Description  
(e.g., block diagrams)**



*| Architecture Implementation*

**Logic Circuit Description  
(Circuit Schematic Diagrams)**



# Assembly Language

- Basic job of a CPU: execute lots of instructions
- Instructions are the primitive operations that the CPU may execute

Like a sentence: *operations* (verbs) applied to *operands* (objects), processed in sequence ...

- Different CPUs implement different sets of instructions. The set of instructions a particular CPU implements is an Instruction Set Architecture (ISA)

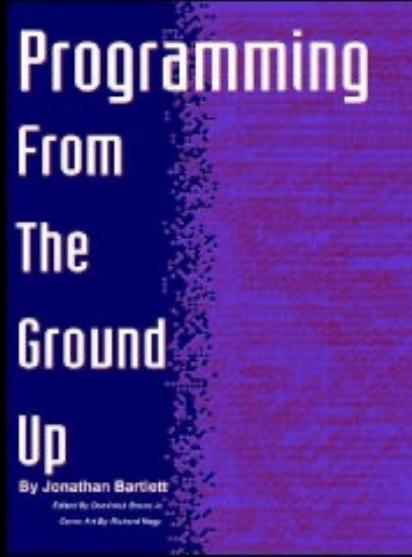
Examples: ARM (cell phones), Intel x86 (i9, i7, i5, i3), IBM/Motorola PowerPC (old Macs), MIPS, RISC-V, ...



# Book: Programming From the Ground Up

“A new book was just released which is based on a new concept – teaching computer science through assembly language (Linux x86 assembly language, to be exact). This book teaches how the machine itself operates, rather than just the language. I've found that the key difference between mediocre and excellent programmers is whether or not they know assembly language. Those that do tend to understand computers themselves at a much deeper level. Although [almost!] unheard of today, this concept isn't really all that new -- there used to not be much choice in years past. Apple computers came with only BASIC and assembly language, and there were books available on assembly language for kids. This is why the old-timers are often viewed as 'wizards': they had to know assembly language programming.”

-- *slashdot.org comment, 2004-02-05*



# Instruction Set Architectures

- Early trend was to add more and more instructions to new CPUs to do elaborate operations

VAX architecture had an instruction to multiply polynomials!

- Reduced Instruction Set Computer (RISC) philosophy:

Cocke IBM, Patterson, Hennessy, 1980s.

Keep the instruction set small and simple, makes it easier to build *fast* hardware

Let software do complicated operations by composing simpler ones

This went against the conventional wisdom of the time  
(he who laughs last, laughs best)

# Patterson and Hennessy win Turing!





# RISC-V Architecture

# IBM 360 Green Card

- New open-source, license-free ISA spec
    - Supported by growing shared software ecosystem
    - Appropriate for all levels of computing system, from microcontrollers to supercomputers
    - 32-bit, 64-bit, and 128-bit variants  
(class/textbook uses 32-bit)

## ■ Why RISC-V instead of Intel x86-64?

- RISC-V is simple and elegant.  
We don't need to get bogged down in gritty details

- RISC-V has exponential adoption, from microcontrollers to CPUs to warehouse-scale computers

<https://cs61c.org/> → Resources

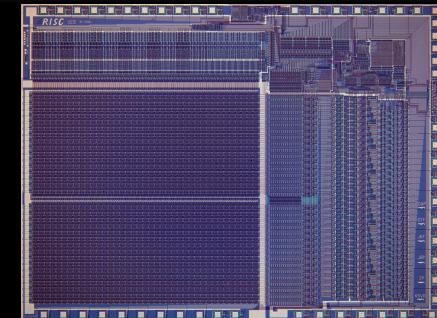
## 07 Intro to RISC-V (8)

# RISC-V Green Card

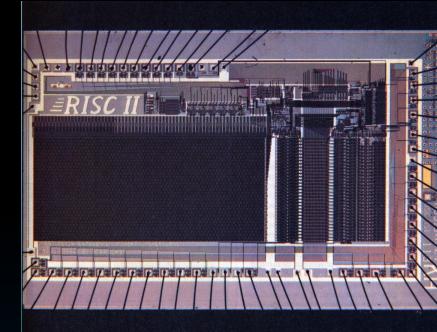
Garcia, Kao

# RISC-V Origins

- Started in Summer 2010 to support open research and teaching at UC Berkeley
  - Lineage can be traced to RISC-I/II projects (1980s)
- As the project matured, it migrated to RISC-V foundation ([www.riscv.org](http://www.riscv.org))
- Many commercial and research projects based on RISC-V
  - Open-source and proprietary
  - Widely used in education
- Read more:
  - <https://riscv.org/risc-v-history/>
  - <https://riscv.org/risc-v-genealogy/>



RISC-I



RISC-II



# Elements of Architecture: Registers

- The RISC-V Instruction Set Architecture
- Elements of Architecture: Registers
- Add/Sub Instructions
- Immediates

Preliminary discussion of the logical design of an electronic computing instrument<sup>1</sup>

Arthur W. Burks / Herman H. Goldstine /  
John von Neumann

“instruction sets”

3.1. It is easy to see by formal-logical methods that there exist codes that are *in abstracto* adequate to control and cause the execution of any sequence of operations which are individually available in the machine and which are, in their entirety, conceivable by the problem planner. The really decisive considerations from the present point of view, in selecting a code, are more of a practical nature: simplicity of the equipment demanded by the code, and the clarity of its application to the actually important problems together with the speed of its handling of those problems. It would take us much too far afield to discuss these questions at all generally or from first principles. We will therefore restrict ourselves to analyzing only the type of code which we now envisage for our machine.

- Instruction set for a particular architecture (e.g. RISC-V) is represented by the assembly language
- Each line of assembly code represents one instruction for the computer

add                     $x_1, x_2, x_3$   
                        \underbrace{\hspace{1cm}} \quad \underbrace{\hspace{1cm}} \quad \underbrace{\hspace{1cm}}  
                        operation name                    registers

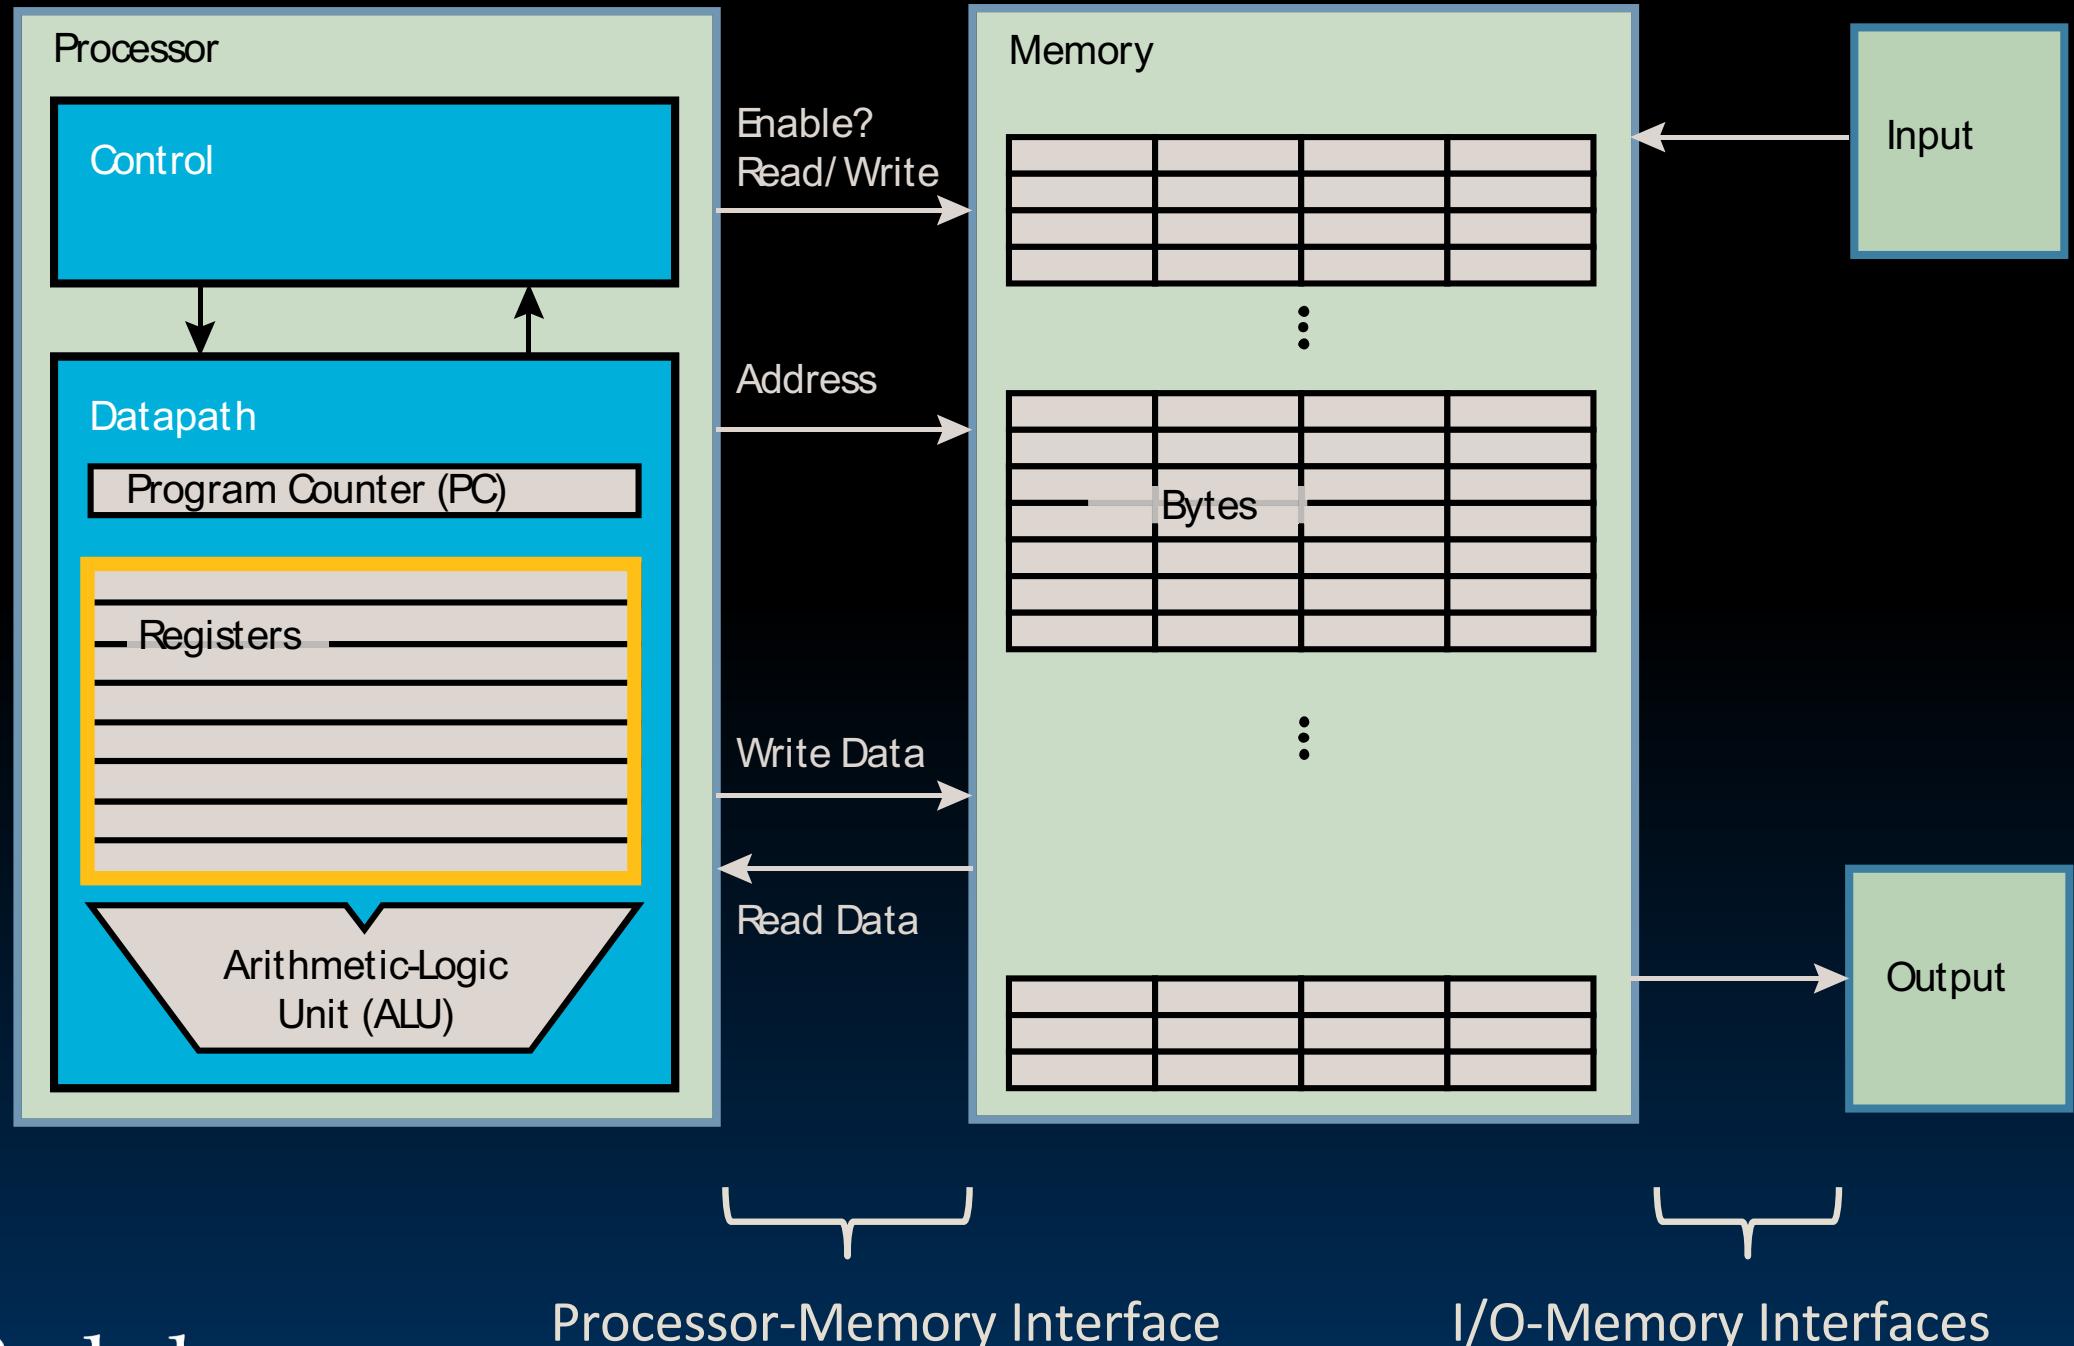
## C, Java

- In C (and most high-level languages), variables declared first and given a type
  - `int fahr, celsius;`  
`char a, b, c, d, e;`
- Each variable can ONLY represent a value of the type it was declared as (cannot mix/match `int/char` variables)
- In high-level languages, variable types determine operation
  - `int *p = ...;`  
`p = p + 2;`  
`int x = 42;`  
`x = 3 * x;`

## RISC-V

- Assembly operands are registers
  - Registers are limited number of special locations, built directly into the hardware
  - RISC-V: Operations can only be performed on register operands
- In assembly language, the registers have no type
  - Register contents are just bits
  - Operation determines “type,” i.e., how register contents are treated
    - E.g., as value, memory address, etc.

# Registers are Inside the Processor

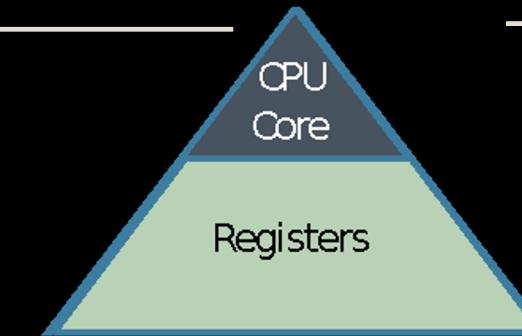


**Benefit:**  
Since registers are directly in hardware, they're very fast (faster than 0.25ns)

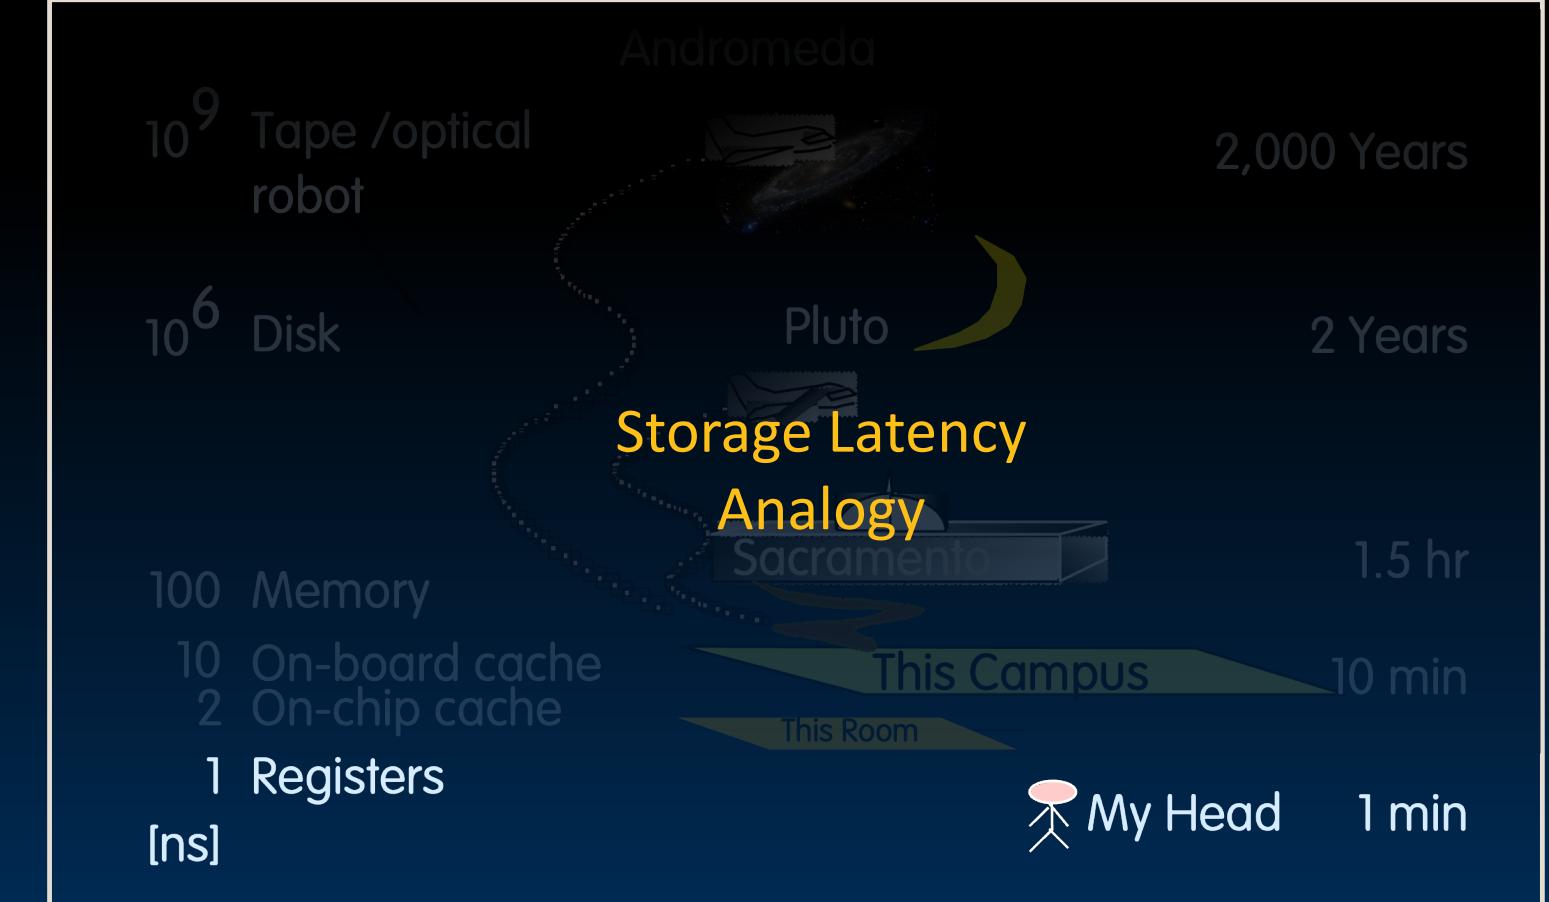
- Speed of light:  
 $3 \times 10^8 \text{ m/s} = 0.3 \text{ m/ns} = 30 \text{ cm/ns} = 10 \text{ cm}/0.3 \text{ ns}!!!$
- 0.3ns is the clock period of a 3.33GHz computer

# Great Idea #3: Principle of Locality / Memory Hierarchy

Processor  
chip



Extremely fast  
Extremely expensive  
Tiny capacity



Garcia, Kao



# 32 Registers in RISC-V

- **Drawback: Each ISA a predetermined number of registers**
  - Why? Registers are built into hardware
  - Solution: RISC-V code must be very carefully put together to efficiently use registers
    - Nowadays: compiler maps C variables to RISC-V registers across declarations, function calls, etc.
- **32 registers in RISC-V**
  - Why 32? Smaller is faster, but too small is bad
  - Goldilocks principle (“This porridge is too hot; This porridge is too cold; this porridge is just right”)
- **Each RISC-V register is 32 bits wide (in RV32 variant)**
  - Groups of 32 bits called a word in RV32

# Register Names and Numbers

- Registers are numbered from 0 to 31
  - Referred to by number **x0 – x31**
- **x0 is special, always holds value zero**
  - So only 31 registers able to hold variable values
- **Each register can be referred to by number or name**
  - We'll explain names next week
- **Ok, enough already...gimme my RV32!!!**

# Add/Sub Instructions

- The RISC-V Instruction Set Architecture
- Elements of Architecture:  
Registers
- Add/Sub Instructions
- Immediates

# Higher-Level Language vs. Assembly Language (2/2)

C, Java

- **Each line can contain multiple operations**

```
a = b * 2 - (arr[2] + *p);
```

- **Common operations are:**

- =, +, -, \*, /
- Here, \* includes both multiply and dereference

RISC-V

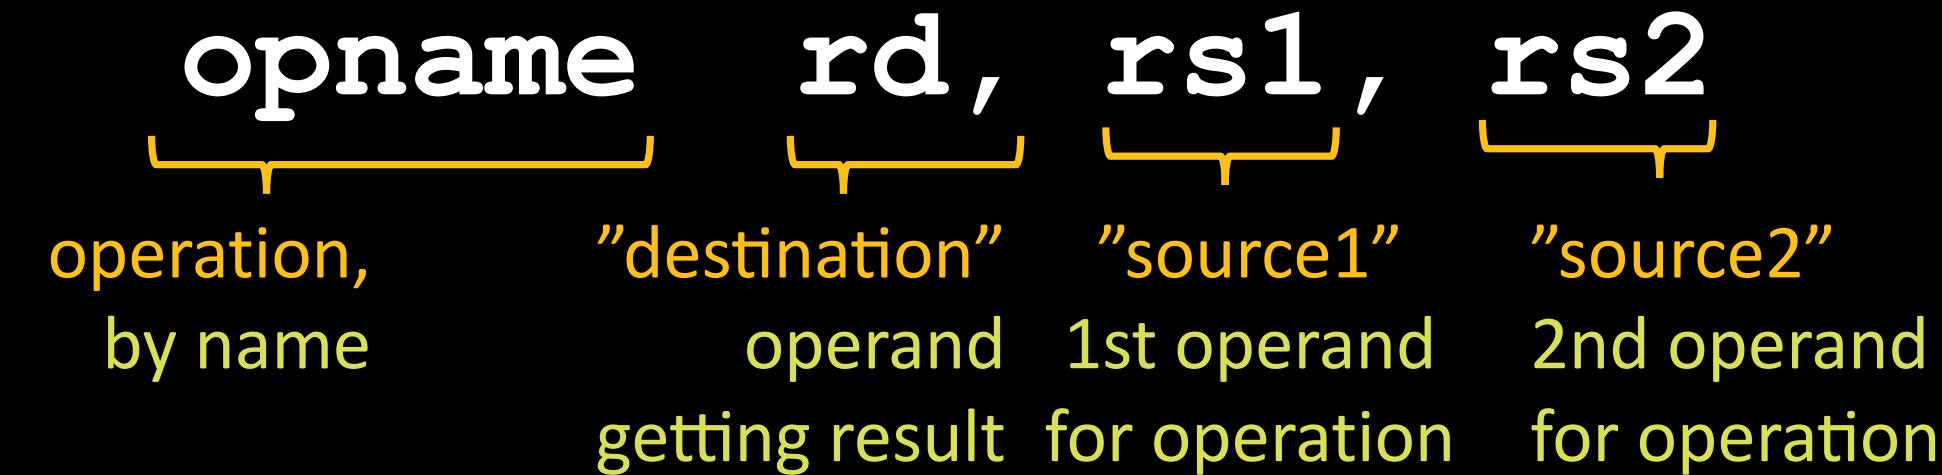
- **Each statement is an instruction**
  - An instruction executes exactly one short list of simple commands
- **Each line of assembly code contains at most 1 instruction**

```
add x1, x2, x3
```

```
sub x1, x1, x4
```

The compiler translates from higher-level language down to assembly...  
...so RISC-V instructions are often closely related to common C/Java operations.

# RISC-V Arithmetic Instruction Syntax



- **Syntax is rigid:**
  - 1 operator, 3 operands
  - Why? Keep hardware simple via regularity ([more in a few weeks](#))

# RISC-V Addition and Subtraction

<b>opname</b>	<b>rd</b> , <b>rs1</b> , <b>rs2</b>
operation, by name	"destination" operand getting result
	"source1" 1st operand for operation
	"source2" 2nd operand for operation

## ▪ Addition

RISC-V    **add x1 ,x2 ,x3**  
C       $a = b + c;$

## ▪ Subtraction

RISC-V    **sub x4 ,x5 ,x6**  
C       $d = e - f;$   
(remember!  
order of operands  
matter in **sub**)

# RISC-V Arithmetic, Example 1

- How could we translate this C statement?

```
a = b + c + d - e;
```

x10 x1 x2 x3 x4

- Assume the above mapping for C variables  $\leftrightarrow$  RISC-V registers.
- Solution: Break into multiple instructions!

```
add x10, x1, x2
add x10, x10, x3
sub x10, x10, x4
```

```
# a_temp = b + c
# a_temp = a_temp + d
# a = a_temp - e
```

A single line of C may break up into several lines of RISC-V.

In-line comments prefixed by #, work like C99's //.  
Unlike C, no multi-line comment /\* \*/ support.

# RISC-V Arithmetic, Example 2

- How do we do this?

$$f = (g + h) - (i + j);$$

x19    x20    x21    x22    x23    x5    x6

- Assume the above mapping for C variables  $\leftrightarrow$  RISC-V registers.

- Solution: Use intermediate temporary registers!

add x5, x20, x21	# a_temp = g + h
add x6, x22, x23	# b_temp = i + j
sub x19, x5, x6	# f = (g + h) - (i + j)

- Note: A good compiler may actually do the following. Why?

$f = g + h - i - j;$     Can use zero temp registers!

# Aside: Apollo Guidance Computer

- Margaret Hamilton  
Director of Software Engineering  
Division of MIT Instrumentation Laboratory
  - Built Apollo Guidance Computer Command Module (navigation, lunar landing)
    - 72KB of memory on board *Eagle* (Apollo 11, 1969)

```

179      TC  BANKCALL    # TEMPORARY, I HOPE HOPE HOPE
180      CADR STOPRATE   # TEMPORARY, I HOPE HOPE HOPE
181      TC  DOWNFLAG   # PERMIT X-AXIS OVERRIDE
  
```

```

245      CAF  CODE500    # ASTRONAUT: PLEASE CRANK THE
246      TC   BANKCALL.  #
247      CADR GOPERF1
248      TCF  GOTOP00H   # TERMINATE
249      TCF  P63SPOT3   # PROCEED SEE IF HE'S LYING
250
251 P63SPOT4   TC  BANKCALL    # ENTER      INITIALIZE LANDING RADAR
252      CADR SETPOS1
253
254      TC   POSTJUMP   # OFF TO SEE THE WIZARD ...
255      CADR BURNBABY
  
```

39 ## It traces back to 1965 and the Los Angeles riots, and was inspired  
 40 ## by disc jockey extraordinaire and radio station owner Magnificent Montague.  
 41 ## Magnificent Montague used the phrase "Burn, baby! BURN!" when spinning the  
 42 ## hottest new records. Magnificent Montague was the charismatic voice of  
 43 ## soul music in Chicago, New York, and Los Angeles from the mid-1950s to  
 44 ## the mid-1960s.



Margaret Hamilton and  
Apollo code  
Wikimedia Commons



Presidential  
Medal of  
Freedom  
(2016)

<https://abcnews.go.com/Technology/apollo-11s-source-code-tons-easter-eggs-including/story?id=40515222>

# Immediates

- The RISC-V Instruction Set Architecture
- Elements of Architecture: Registers
- Add/Sub Instructions
- Immediates

# Immediates

- **Immediates are numerical constants.**
  - They appear often in code; hence, they have separate instructions.
- **Add Immediate instruction:**

RISC-V      **addi x3, x4, 10**

C       $f = g + 10;$

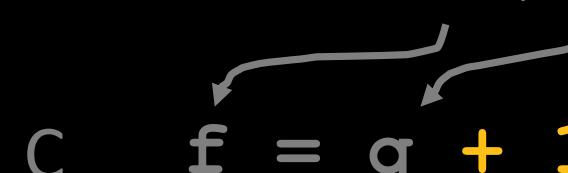
Syntax similar to **add** instruction. But the *last operand must be a number, not a register.*

# Immediates

- Immediates are numerical constants.

They appear often in code; hence, they have separate instructions.

- Add Immediate instruction:

RISC-V      **addi** `x3, x4, 10`  
  
C      `f = g + 10;`

Syntax similar to **add** instruction. But the *last operand* must be *a number*, not a register.

- There is no Subtract Immediate instruction in RISC-V. Why?

While there are **add** and **sub** instructions, to subtract an immediate, just use **addi**:

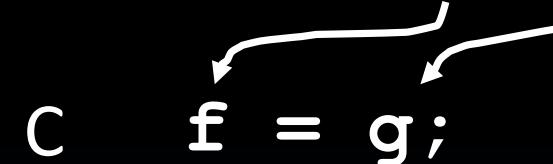
RISC-V      **addi** `x3, x4, -10`  
  
C      `f = g - 10;`

**RISC Philosophy: Reduce** the possible types of operations to an absolute minimum.  
If an operation can be decomposed into a simpler operation, don't include it in the ISA.

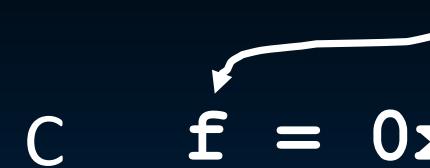
# Register Zero

- One particular immediate, the number zero (**0**), appears very often in code.
- RISC-V hardwires the register zero (**x0**) to value **0**:

RISC-V    add    **x3 , x4 , x0**  
C      f = g;



RISC-V    addi    **x3 , x0 , 0xff**  
C      f = 0xff;



- Defined in hardware, so an instruction    add    **x0 , x3 , x4** will not do anything!

# Concept Check: True or False?

1. Types are associated with declaration in C (normally), but are associated with instructions (operators) in RISC-V.
2. Since there are only 32 registers, we can't write RISC-V for C expressions that contain > 32 variables.
3. If p (stored in x9) were a pointer to an array of ints, then p++ ; would be  
`addi x9 x9 1`

123

- A. FFF
- B. FFT
- C. FTF
- D. FTT
- E. TFF
- F. TFT
- G. TTF
- H. TTT

## L07 Three questions

1. Types are associated with declaration in C (normally), but are associated with instructions (operators) in RISC-V.
2. Since there are only 32 registers, we can't write RISC-V for C expressions that contain > 32 variables.
3. If p (stored in x9) were a pointer to an array of ints, then p++; would be  
`addi x9 x9 1`



# Concept Check: True or False?

1. Types are associated with declaration in C (normally), but are associated with instructions (operators) in RISC-V.
2. Since there are only 32 registers, we can't write RISC-V for C expressions that contain > 32 variables.
3. If p (stored in x9) were a pointer to an array of ints, then p++ ; would be

addi x9 x9 1

1. True, we saw that on an earlier slide.
  2. False. we saw how to break up longer equations to smaller ones already.
  3. False. Don't forget that ints are (usually) 4 bytes wide, so instruction would be
- addi x9, x9, 4.

123

- A. FFF
- B. FFT
- C. FTF
- D. FTT
- E. TFF
- F. TFT
- G. TTF
- H. TTT

# In Conclusion...

- In RISC-V Assembly Language:
  - Registers replace C variables
  - One instruction (simple operation) per line
  - Simpler is Better, Smaller is Faster
- In RV32, words are 32b
- Instructions:  
**add, addi, sub**
- Registers:
  - 32 registers, referred to as **x0 – x31**
  - Zero: **x0**