




UC Berkeley
Teaching Professor
Dan Garcia

CS61C

Great Ideas
in
Computer Architecture
(a.k.a. Machine Structures)

RISC-V Data Transfer

Storing Data in Memory



RV32 So Far...

Addition/subtraction

add rd, rs1, rs2

$R[rd] = R[rs1] + R[rs2]$

sub rd, rs1, rs2

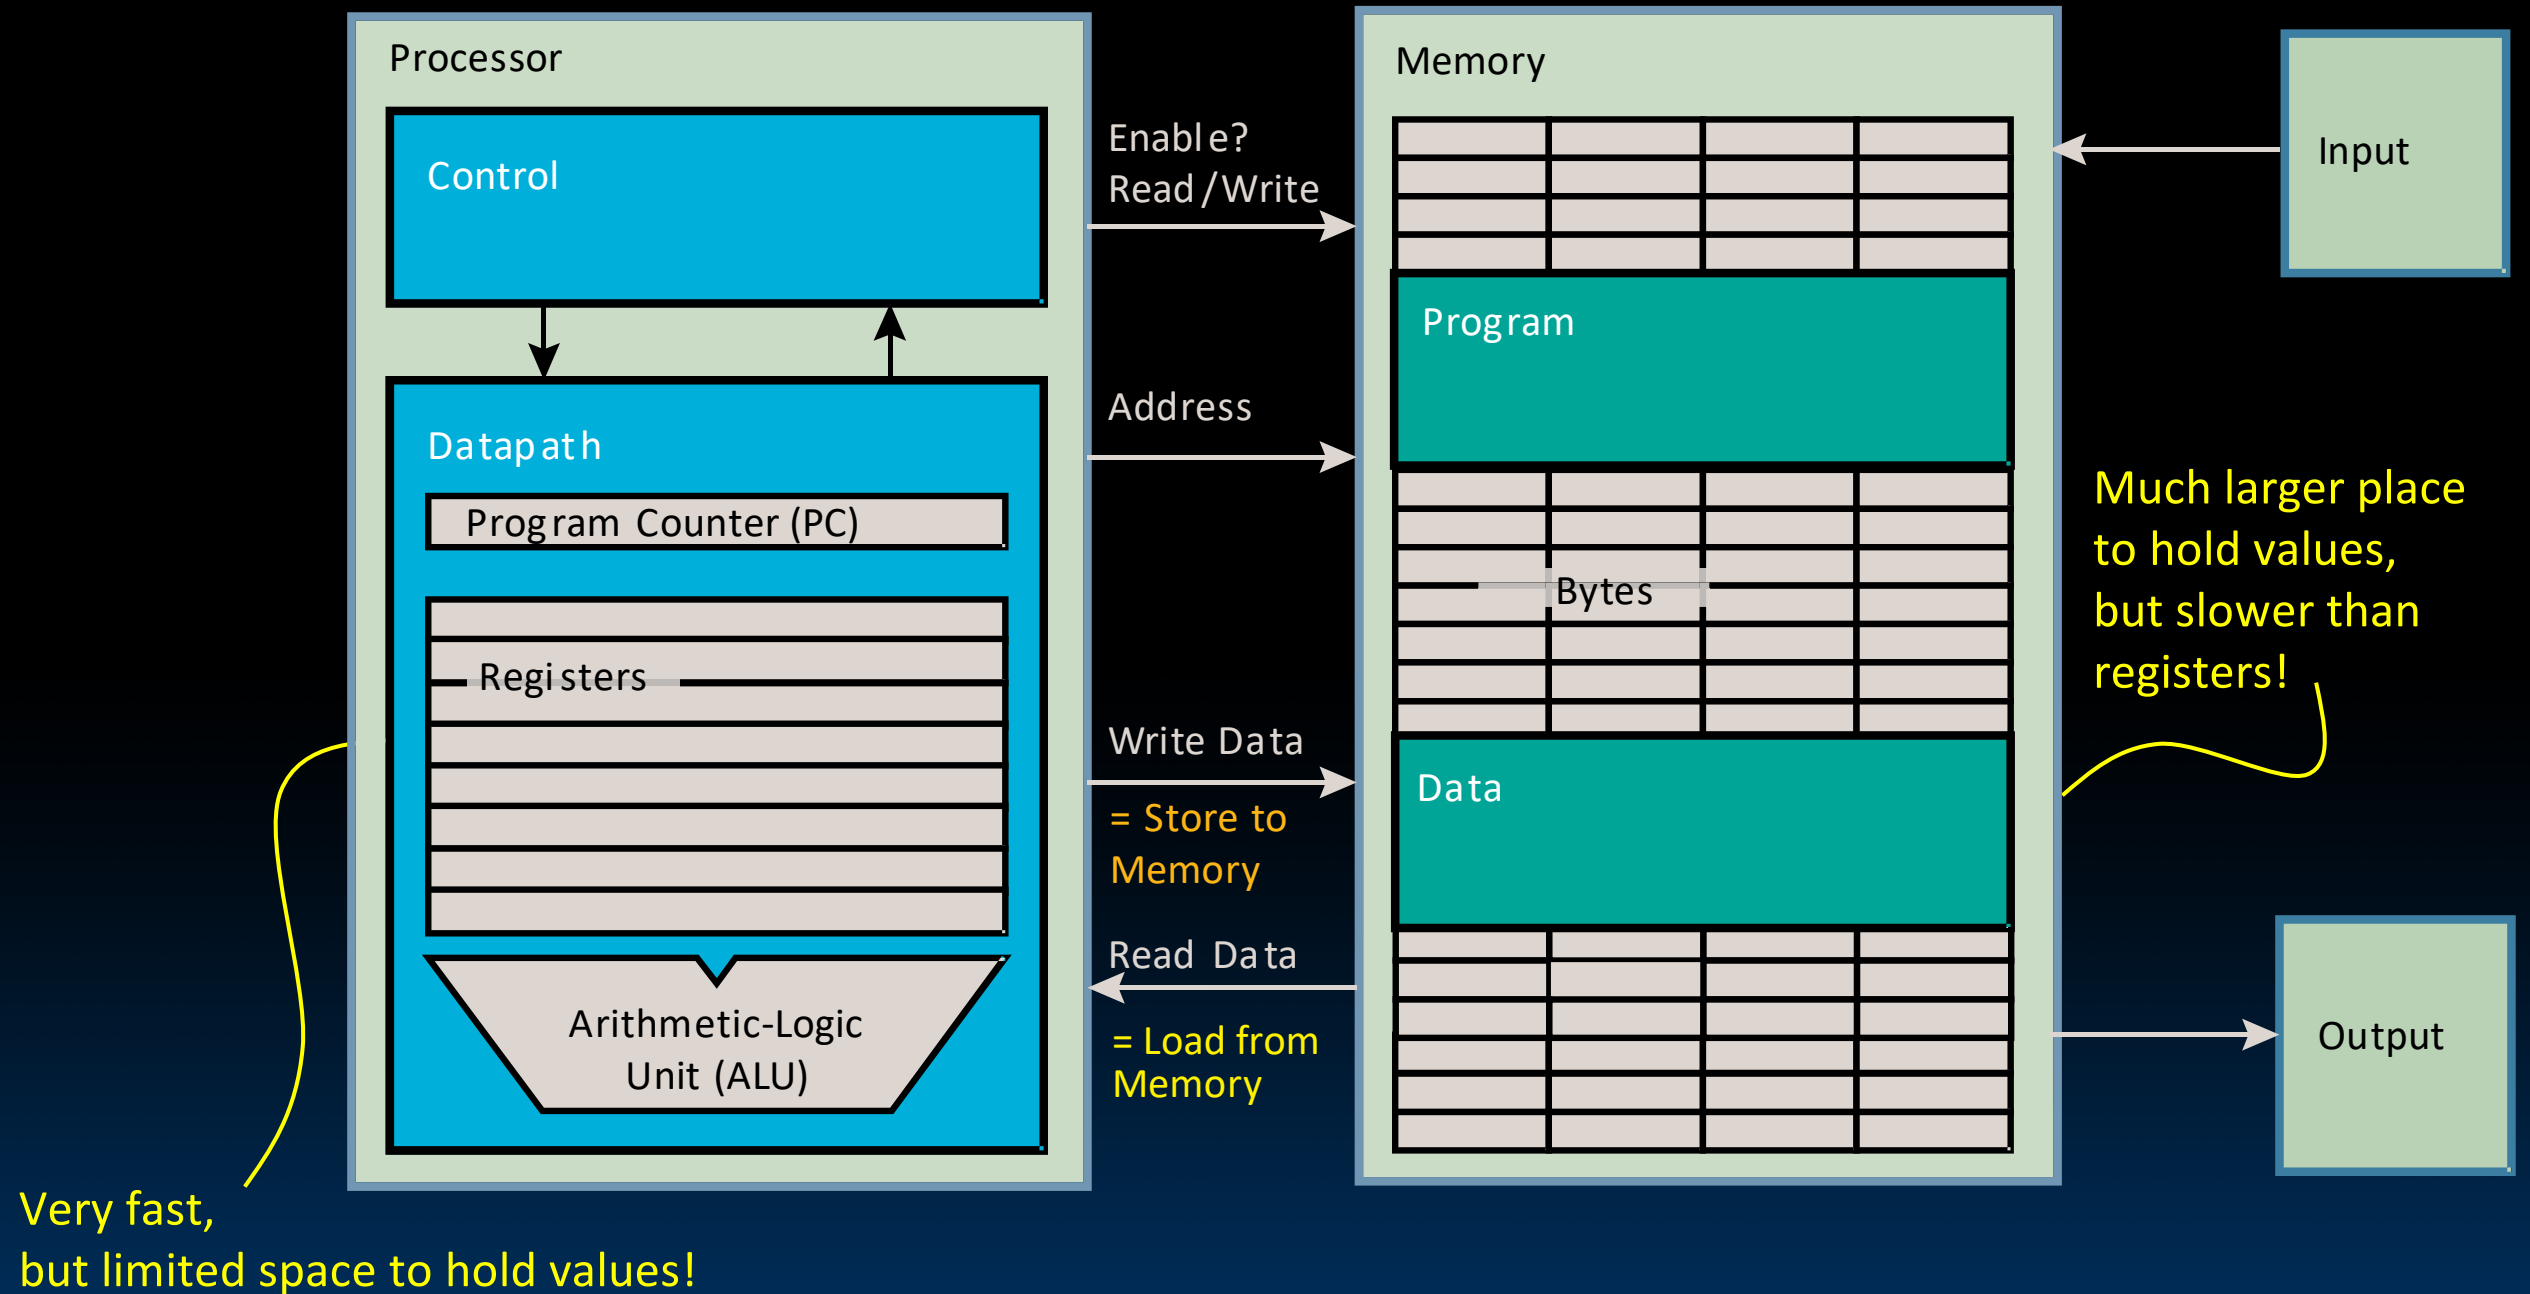
$R[rd] = R[rs1] - R[rs2]$

Add immediate

addi rd, rs1, imm

$R[rd] = R[rs1] + imm$

Data Transfer: **Load from** and **Store to** memory





Memory Addresses are in Bytes


Data typically smaller than 32 bits, but rarely smaller than 8 bits (e.g., char type)—works fine if everything is a multiple of 8 bits


8 bit chunk is called a *byte* (1 word = 4 bytes)

Memory addresses are really in *bytes*, not words

Word addresses are 4 bytes apart

- Word address is same as address of rightmost byte – least-significant byte (i.e. Little-endian convention)





Memory Addresses are in Bytes

Data typically smaller than 32 bits, but rarely smaller than 8 bits (e.g., char type)—works fine if everything is a multiple of 8 bits

8 bit chunk is called a **byte** (1 word = 4 bytes)

Memory addresses are really in **bytes**, not words

Word addresses are 4 bytes apart

- Word address is same as address of rightmost byte – least-significant byte (i.e. Little-endian convention)

Least-significant byte
in a word ↓

15	14	13	12				
11	10	9	8				
7	6	5	4				
3	2	1	0				
31	24	23	16	15	8	7	0

Least-significant byte
gets the smallest address



Big Endian vs. Little Endian

The adjective endian has its origin in the writings of 18th century writer Jonathan Swift. In the 1726 novel Gulliver's Travels, he portrays the conflict between sects of Lilliputians divided into those breaking the shell of a boiled egg from the big end or from the little end. He called them the "Big-Endians" and the "Little-Endians".

- The order in which BYTES are stored in memory
- Bits always stored as usual within a byte (E.g., $0xC2 = 0b\ 1100\ 0010$)

Consider the number 1025 as we typically write it:

BYTE3	BYTE2	BYTE1	BYTE0
00000000	00000000	0000100	00000001

Big Endian

ADDR3	ADDR2	ADDR1	ADDR0
BYTE0	BYTE1	BYTE2	BYTE3
00000001	0000100	00000000	00000000

Examples

Names in China or Hungary (e.g., Garcia Dan)

Java Packages: (e.g., org.mypackage.HelloWorld)

Dates in ISO 8601 YYYY-MM-DD (e.g., 2020-09-07)

Eating Pizza crust first

Little Endian

ADDR3	ADDR2	ADDR1	ADDR0
BYTE3	BYTE2	BYTE1	BYTE0
00000000	00000000	0000100	00000001

Examples

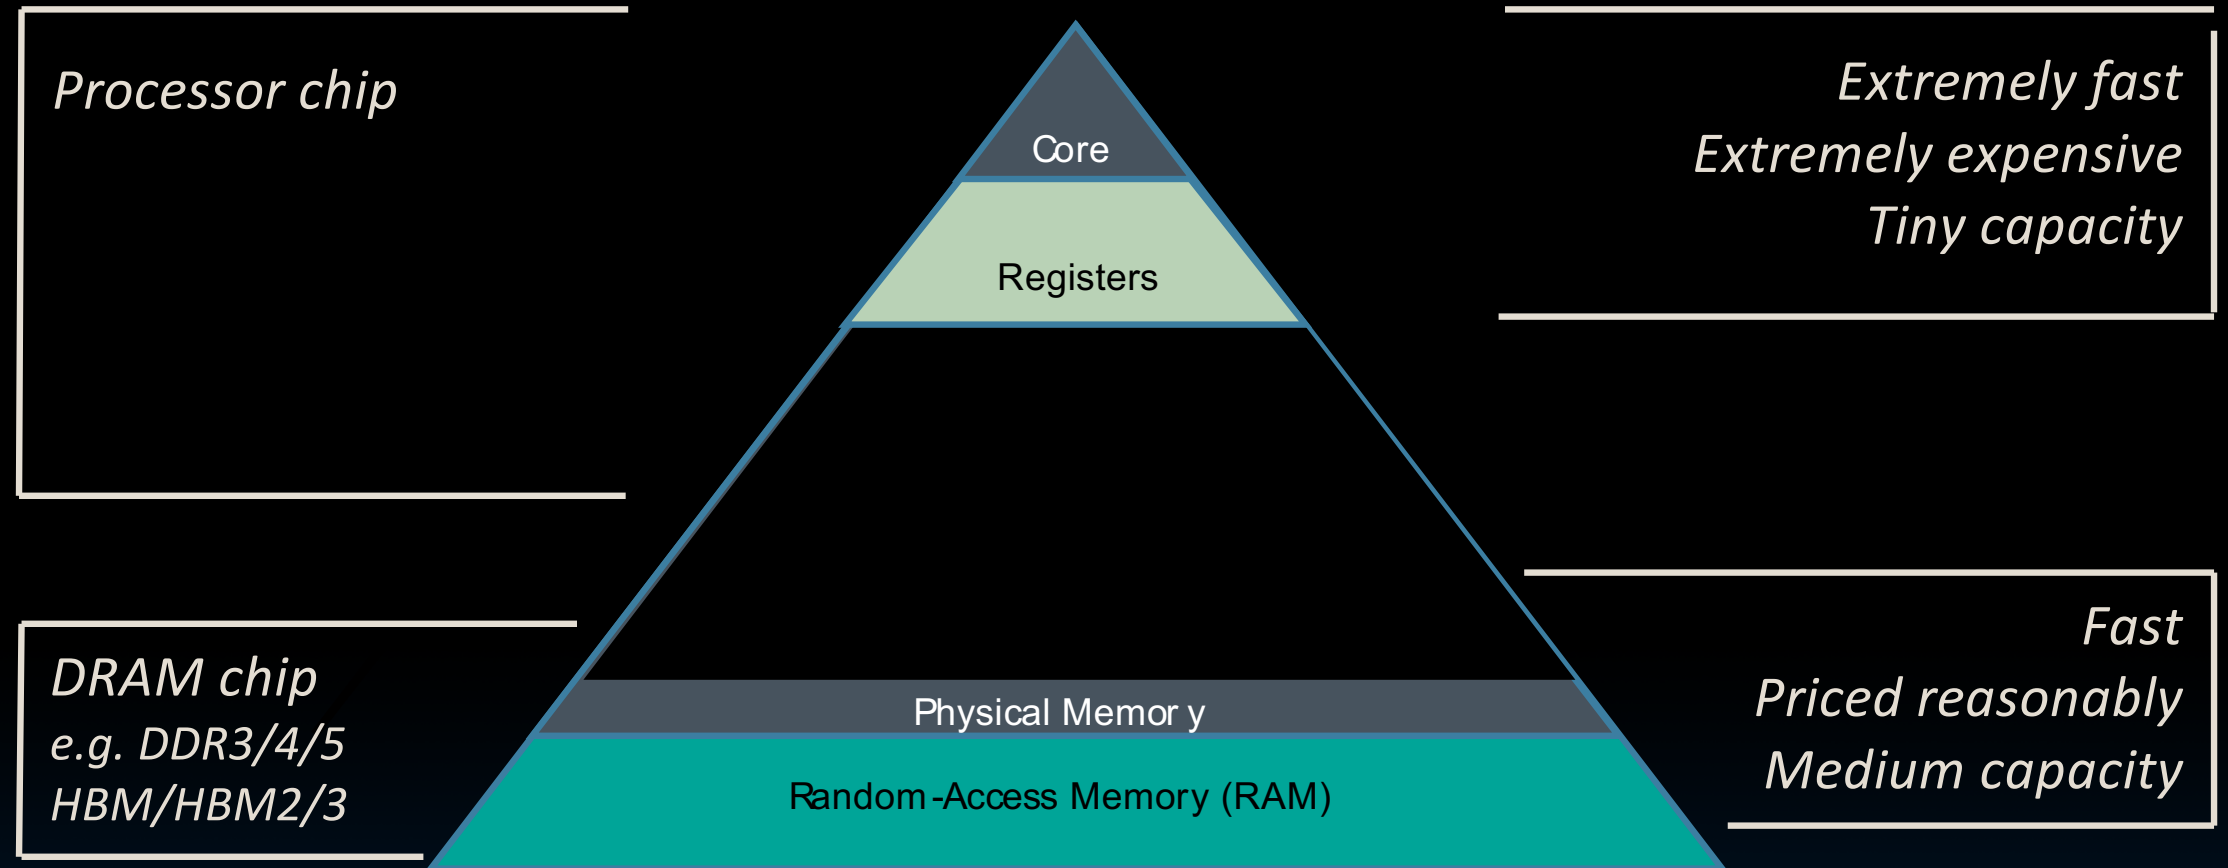
Names in the US (e.g., Dan Garcia)

Internet names (e.g., cs.berkeley.edu)

Dates written in Europe DD/MM/YYYY (e.g., 07/09/2020)

Eating Pizza skinny part first


Great Idea #3: Principle of Locality / Memory Hierarchy






Speed of Registers vs. Memory


- Given that
 - Registers: 32 words (128 Bytes)
 - Memory (DRAM): Billions of bytes (2 GB to 96 GB on laptop)
- and physics dictates...
 - Smaller is faster
- How much faster are registers than DRAM??
 - About 50-500 times faster!
(in terms of latency of one access - tens of ns)
 - But subsequent words come every few ns



Jim Gray's Storage Latency Analogy: How Far Away is the Data?



Jim Gray
Turing Award
B.S. Cal
1966
Ph.D. Cal
1969





Load from Memory to Register

C code

```
int  A[100];  
g = h + A[3];
```



Using Load Word (`lw`) in RISC-V:

```
lw  x10, 12 (x15) # Reg x10 gets A[3]  
add x11, x12, x10 # g = h + A[3]
```

Note: `x15` – base register (pointer to `A[0]`)

`12` – offset in bytes

Offset must be a constant known at assembly time



Store from Register to Memory

C code

```
int  A[100];  
A[10] = h + A[3];
```

Using Store Word (**sw**) in RISC-V:

```
lw  x10, 12(x15)  # Temp reg x10 gets A[3]  
add x10, x12, x10  # Temp reg x10 gets h + A[3]  
sw  x10, 40(x15)  # A[10] = h + A[3]
```



Data flow

Note: **x15** – base register (pointer)

12, 40 – offsets in bytes

x15+12 and **x15+40** must be multiples of 4



Loading and Storing Bytes

- In addition to word data transfers (lw, sw), RISC-V has byte data transfers:

- load byte: **lb**
- store byte: **sb**

- Same format as **lw**, **sw**

- E.g., **lb x10, 3(x11)**

- contents of memory location with address = sum of “3” + contents of register **x11** is copied to the low byte position of register **x10**.

RISC-V also has “unsigned byte” loads (**lbu**) which zero extends to fill register. Why no unsigned store byte ‘**sbu**’?

x10:

xxxxx xxxxx xxxxx xxxxx xxxxx xxxxx

xzzz zzzz

...is copied to “sign-extend”

This bit

byte loaded



What ends up in x12 ?

```
addi x11,x0,0x39C
```

```
sw x11,0(x5)
```

```
lb x12,0(x5)
```

x5

x11
x12

Memory



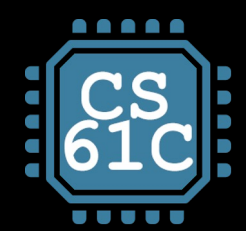


Example: Translate $*x = *y$

We want to translate $*x = *y$ into RISC-V
 x, y ptrs stored in: $x3$ $x5$

```
1: add x3, x5, zero
2: add x5, x3, zero
3: lw  x3, 0(x5)
4: lw  x5, 0(x3)
5: lw  x8, 0(x5)
6: sw  x8, 0(x3)
7: lw  x5, 0(x8)
8: sw  x3, 0(x8)
```

```
1
2
3
4
5→6
6→5
7→8
```



And in Conclusion...

Memory is **byte**-addressable, but **lw** and **sw** access one **word** at a time.

A pointer (used by **lw** and **sw**) is just a memory address, we can add to it or subtract from it (using offset).

Big- vs Little Endian

- Tip: draw lowest byte on the right

New Instructions:

lw, sw, lb, sb, lbu